



# Database Transaction Isolation and Locking in Java

**Kanstantsin Slisenka, Senior Software Engineer  
EPAM Systems**

Aug 25, 2016

# About me



[kanstantsin\\_slisenka@epam.com](mailto:kanstantsin_slisenka@epam.com)

[github.com/kslisenko](https://github.com/kslisenko)

[www.slisenko.net](http://www.slisenko.net)

## Kanstantsin Slisenka

### EXPERIENCE

- Java backend developer
- Oracle Certified Java 7 Programmer, OCEWCD, OCEJPAD
- Speaker at Java tech talks

### INTERESTED IN

- Java backend, SOA, databases, integration frameworks
- Solution architecture
- High load, fault-tolerant, distributed, scalable systems

**DID YOU USE TX ISOLATION OR  
LOCKING AT YOUR PROJECTS?**

# Agenda

---

- Transaction phenomena and isolation levels
- Pessimistic and optimistic approaches
- Transaction isolation in MySQL
- Database-level locks in MySQL
- JPA features for locking

WHAT IS DATABASE TRANSACTION?

# Database transactions and ACID

- 1 Atomicity
- 2 Consistency
- 3 Isolation
- 4 Durability



# Transaction phenomena

## PROBLEM

- Problem of concurrent updates made by parallel transactions
- No problem if no concurrent updates
- Databases have protection

## PHENOMENA

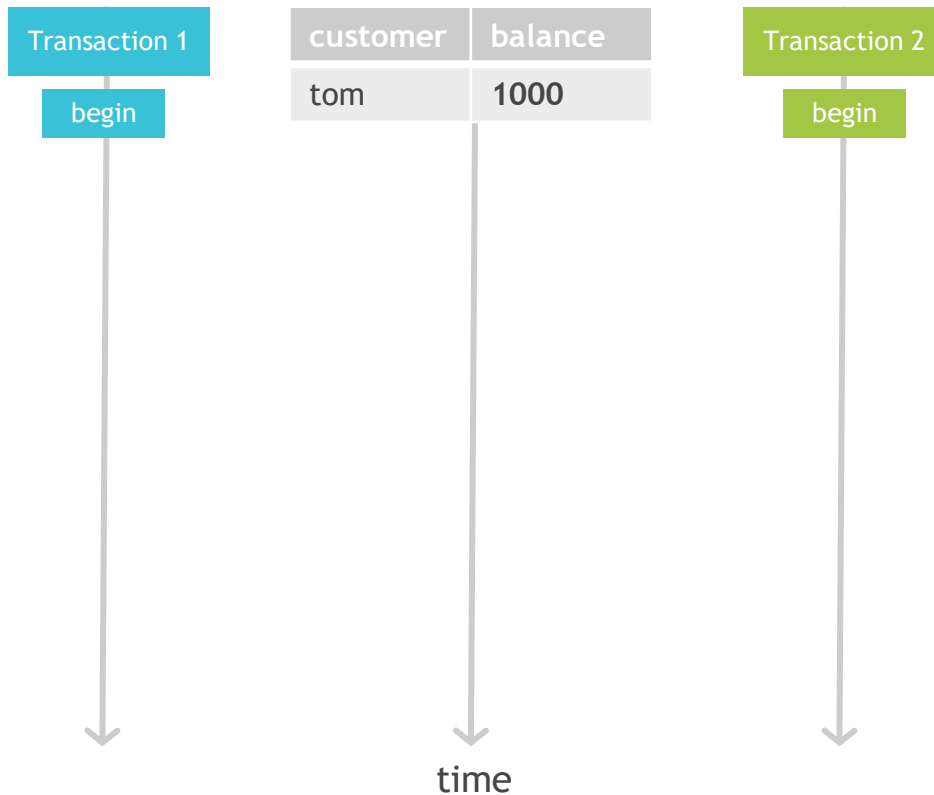
- Dirty read
- Non-repeatable read
- Phantom insert

# Transaction phenomena: dirty read

## PROBLEM

- Transactions can read not committed (dirty) data of each other
- Other transaction rollbacks, decision was made based on never existed data

**DATABASES ARE  
PROTECTED AGAINST  
THIS IN REAL LIFE**



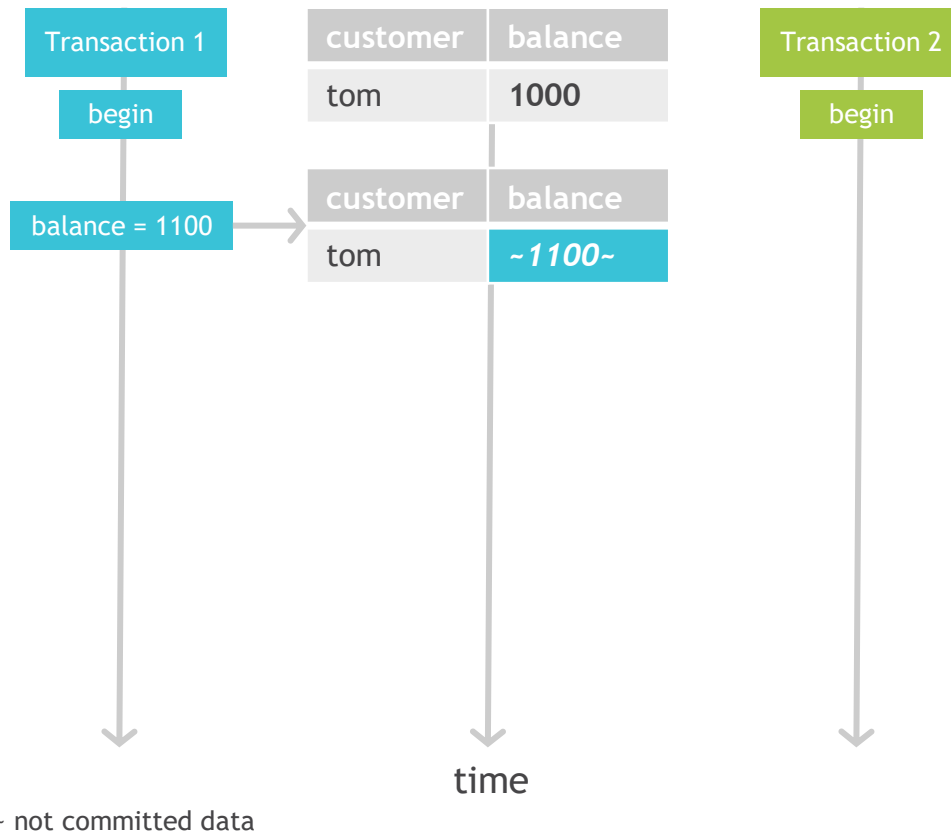


# Transaction phenomena: dirty read

## PROBLEM

- Transactions can read not committed (dirty) data of each other
- Other transaction rollbacks, decision was made based on never existed data

**DATABASES ARE  
PROTECTED AGAINST  
THIS IN REAL LIFE**

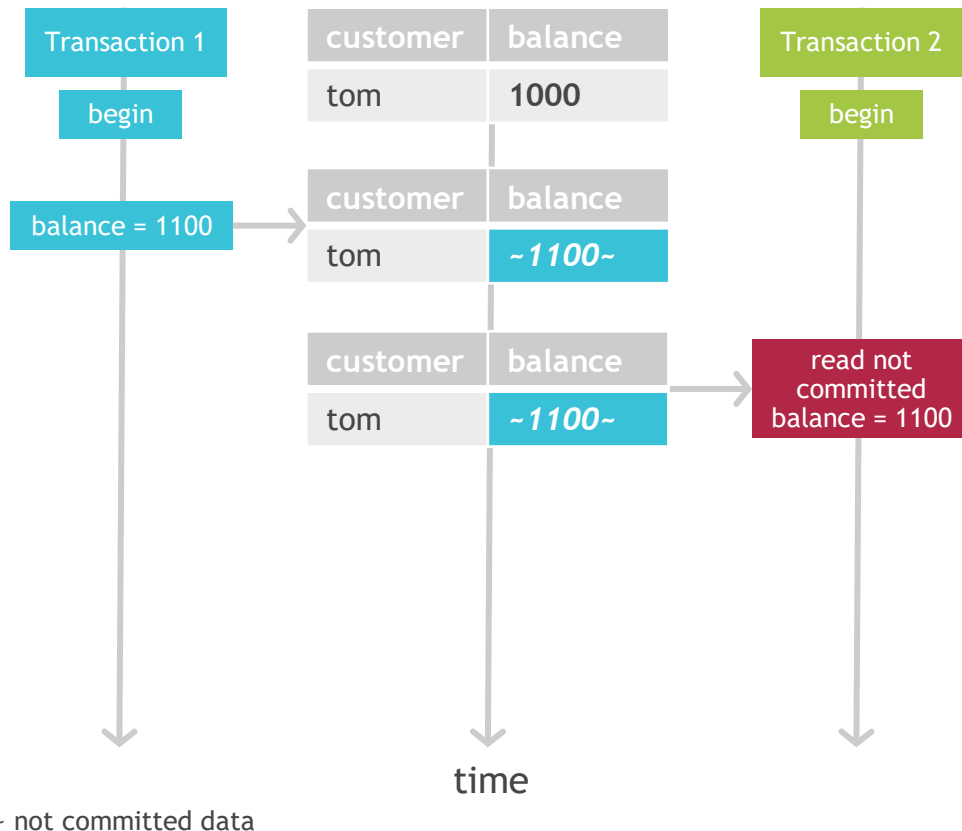


# Transaction phenomena: dirty read

## PROBLEM

- Transactions can read not committed (dirty) data of each other
- Other transaction rollbacks, decision was made based on never existed data

**DATABASES ARE PROTECTED AGAINST THIS IN REAL LIFE**

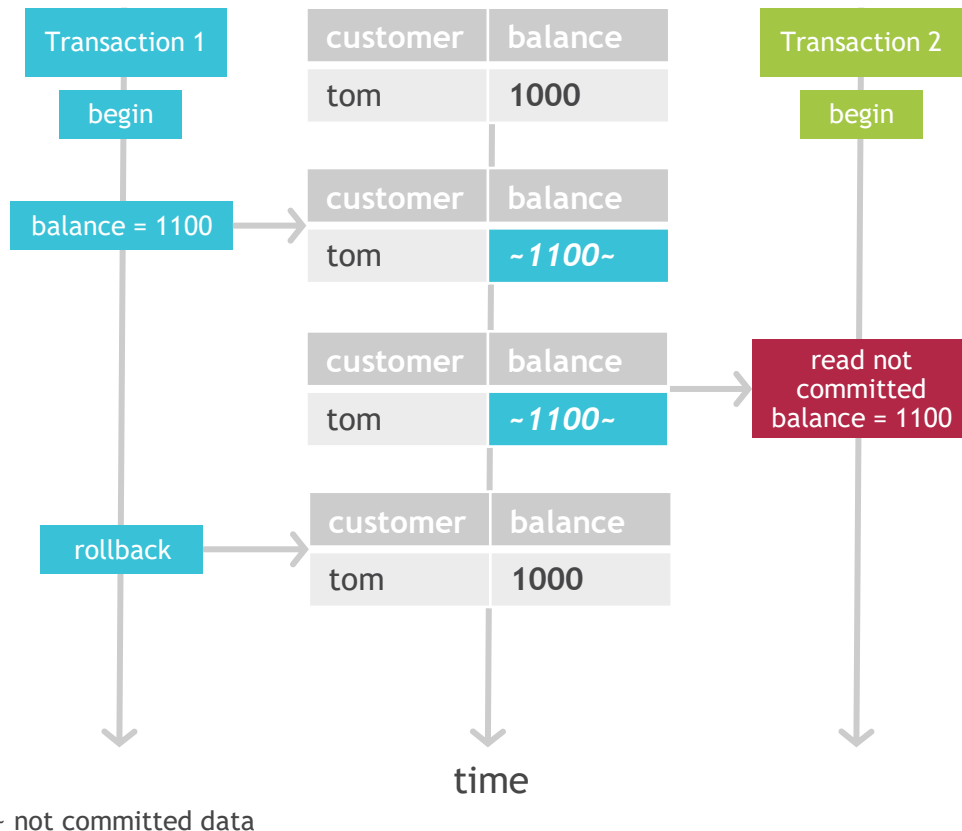


# Transaction phenomena: dirty read

## PROBLEM

- Transactions can read not committed (dirty) data of each other
- Other transaction rollbacks, decision was made based on never existed data

**DATABASES ARE PROTECTED AGAINST THIS IN REAL LIFE**



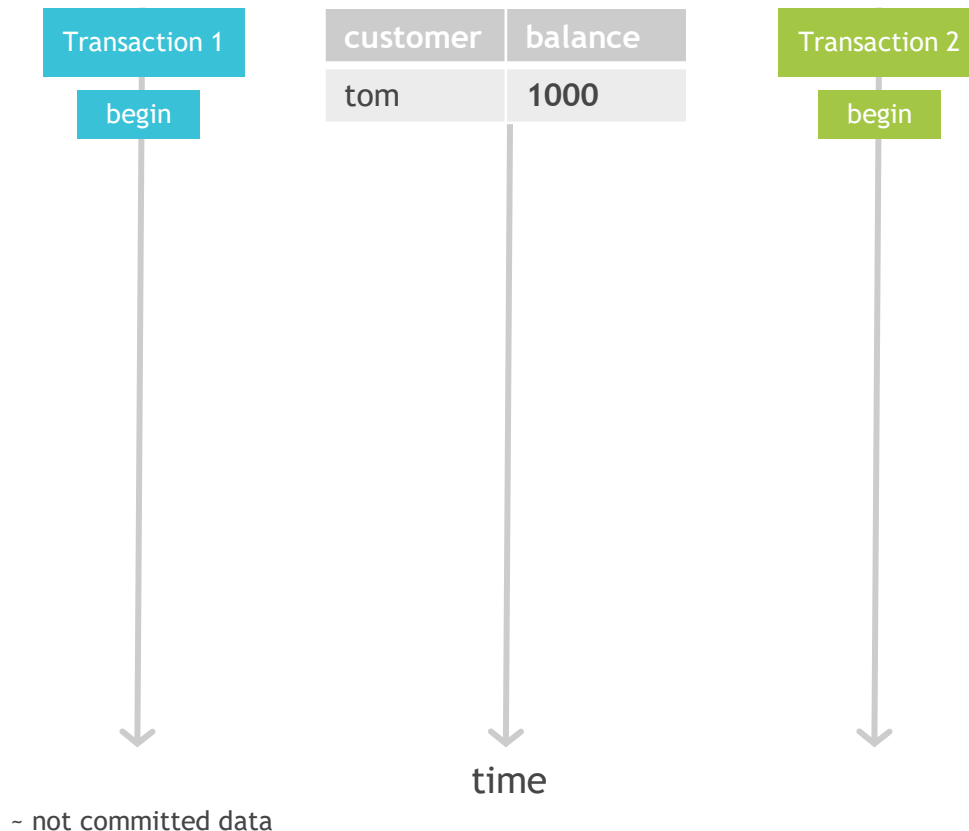
# Transaction phenomena: non-repeatable read

## PROBLEM

- One transaction updates data
- Other transaction reads data several times and get different results

## WHEN WE CAN LIVE WITH THIS

- We are fine with not the most recent data



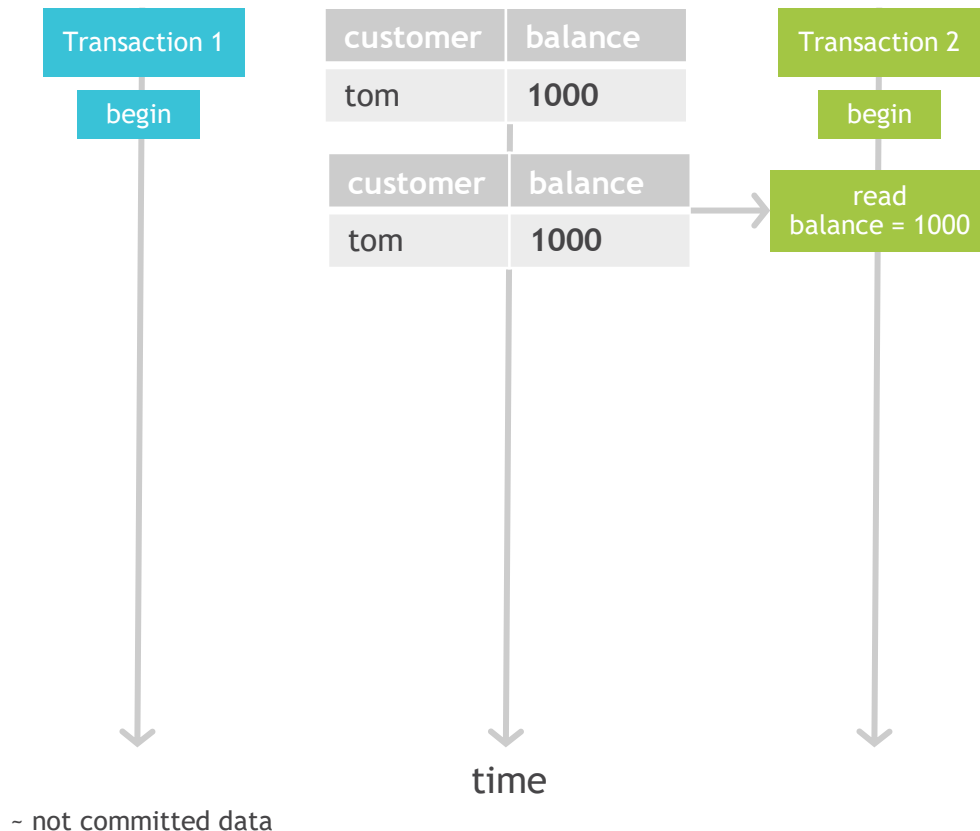
# Transaction phenomena: non-repeatable read

## PROBLEM

- One transaction updates data
- Other transaction reads data several times and get different results

## WHEN WE CAN LIVE WITH THIS

- We are fine with not the most recent data



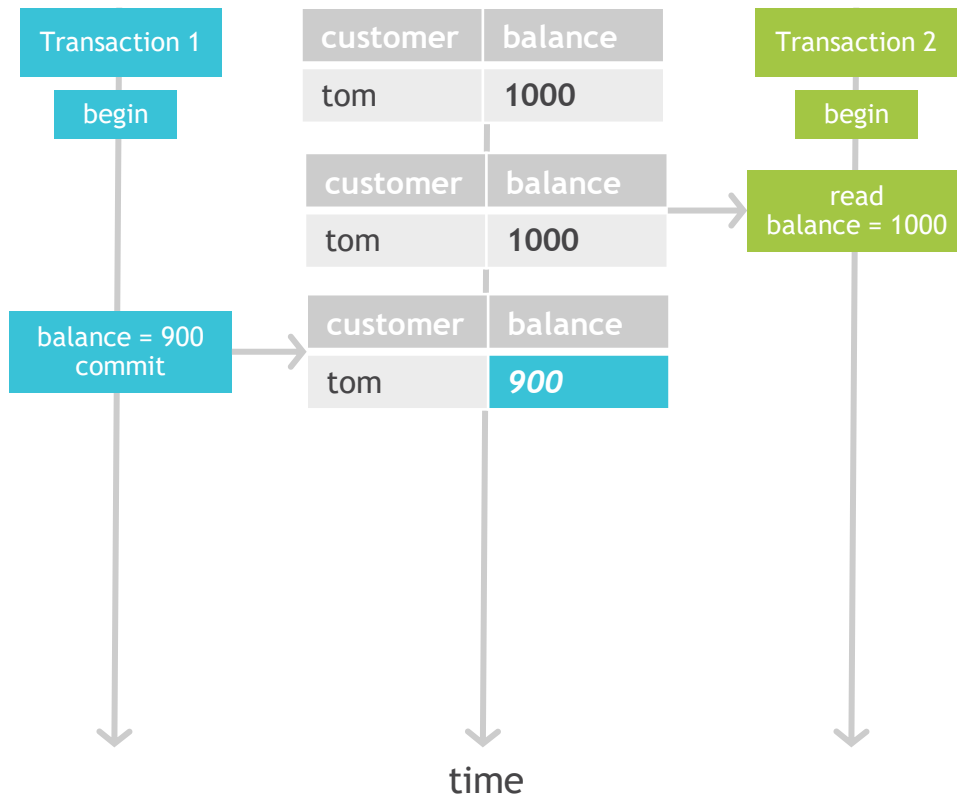
# Transaction phenomena: non-repeatable read

## PROBLEM

- One transaction updates data
- Other transaction reads data several times and get different results

## WHEN WE CAN LIVE WITH THIS

- We are fine with not the most recent data



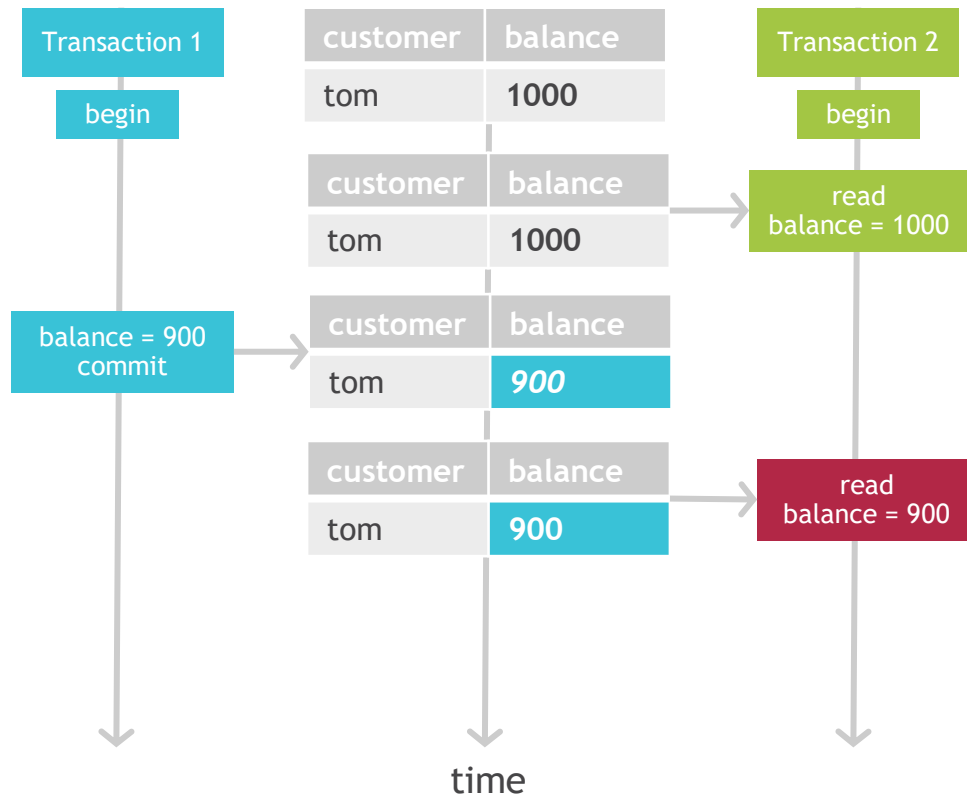
# Transaction phenomena: non-repeatable read

## PROBLEM

- One transaction updates data
- Other transaction reads data several times and get different results

## WHEN WE CAN LIVE WITH THIS

- We are fine with not the most recent data



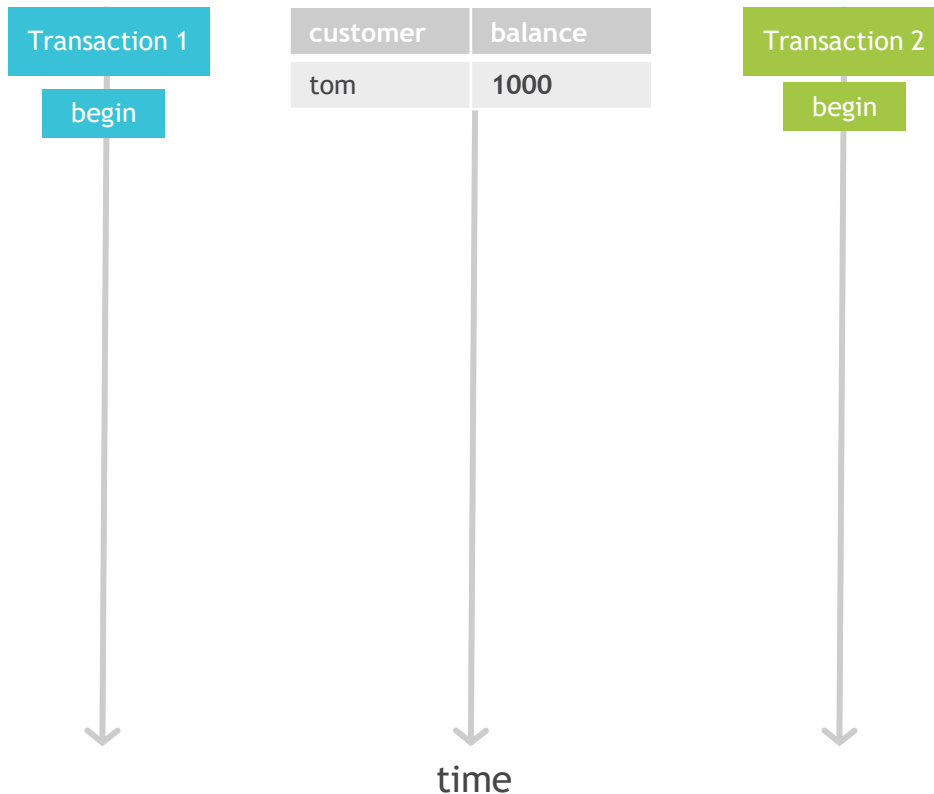
# Transaction phenomena: phantom

## PROBLEM

- One transaction inserts/deletes rows
- Other transaction reads several times and get different number of rows

## WHEN WE CAN LIVE WITH THIS

- Read single rows, not ranges
- We are fine with not the most recent data





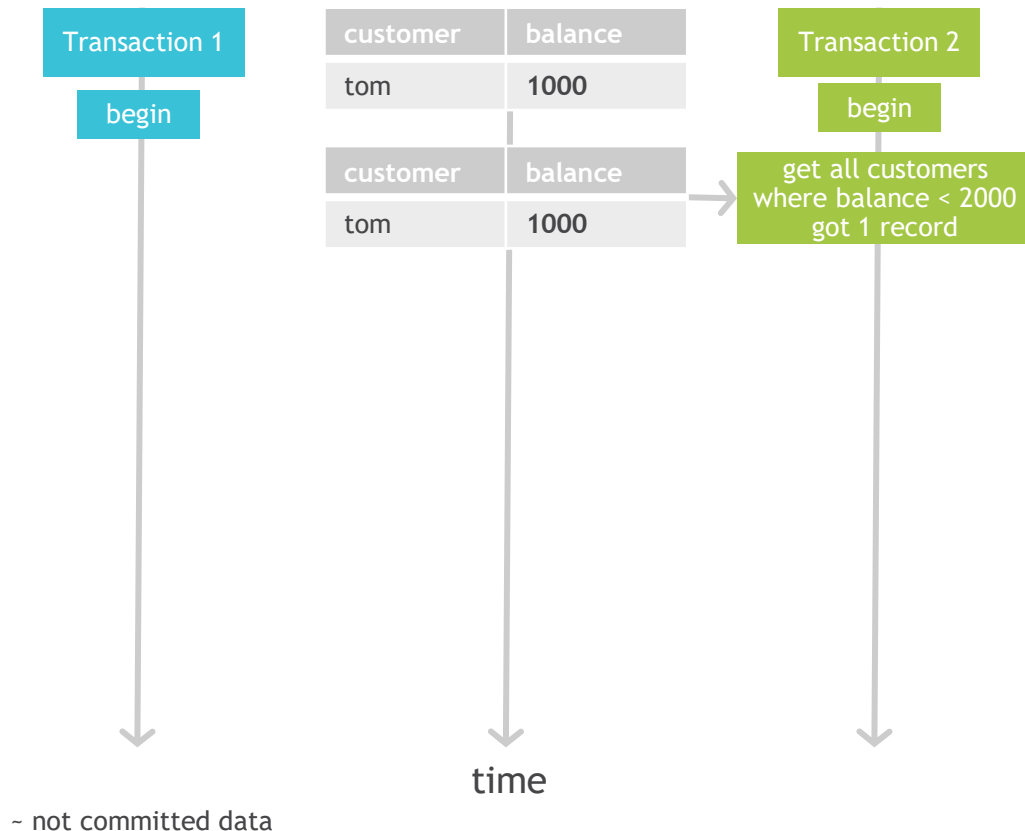
# Transaction phenomena: phantom

## PROBLEM

- One transaction inserts/deletes rows
- Other transaction reads several times and get different number of rows

## WHEN WE CAN LIVE WITH THIS

- Read single rows, not ranges
- We are fine with not the most recent data



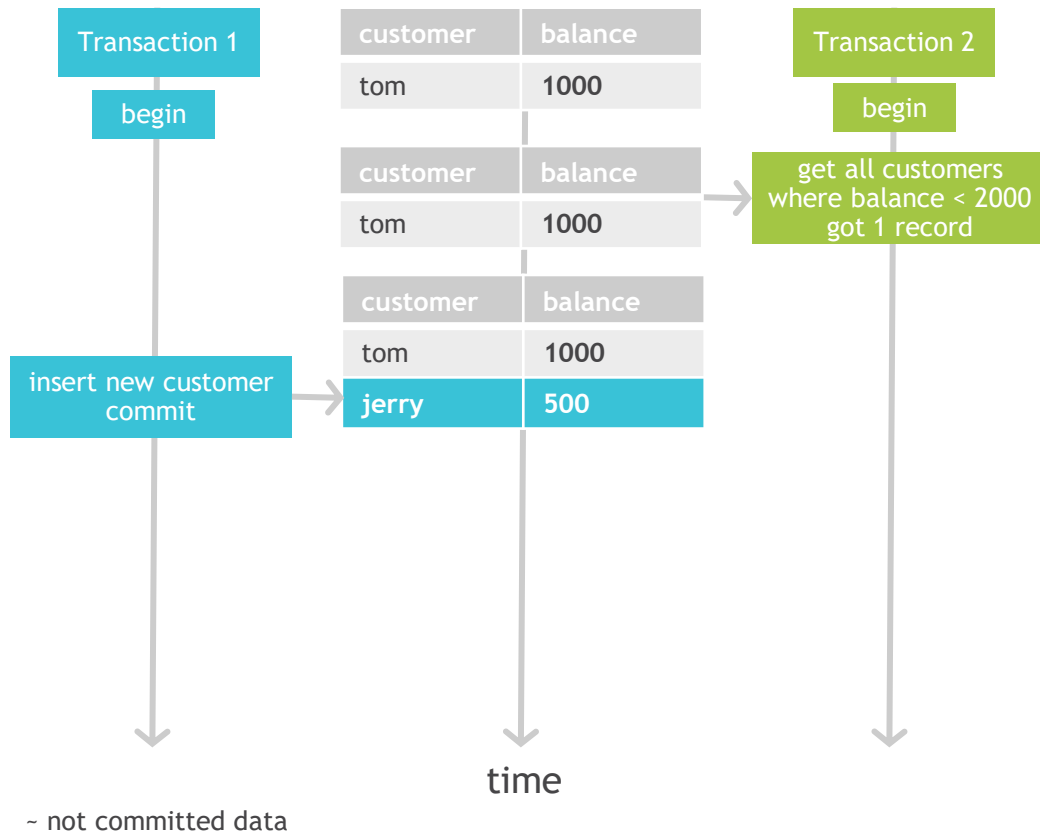
# Transaction phenomena: phantom

## PROBLEM

- One transaction inserts/deletes rows
- Other transaction reads several times and get different number of rows

## WHEN WE CAN LIVE WITH THIS

- Read single rows, not ranges
- We are fine with not the most recent data



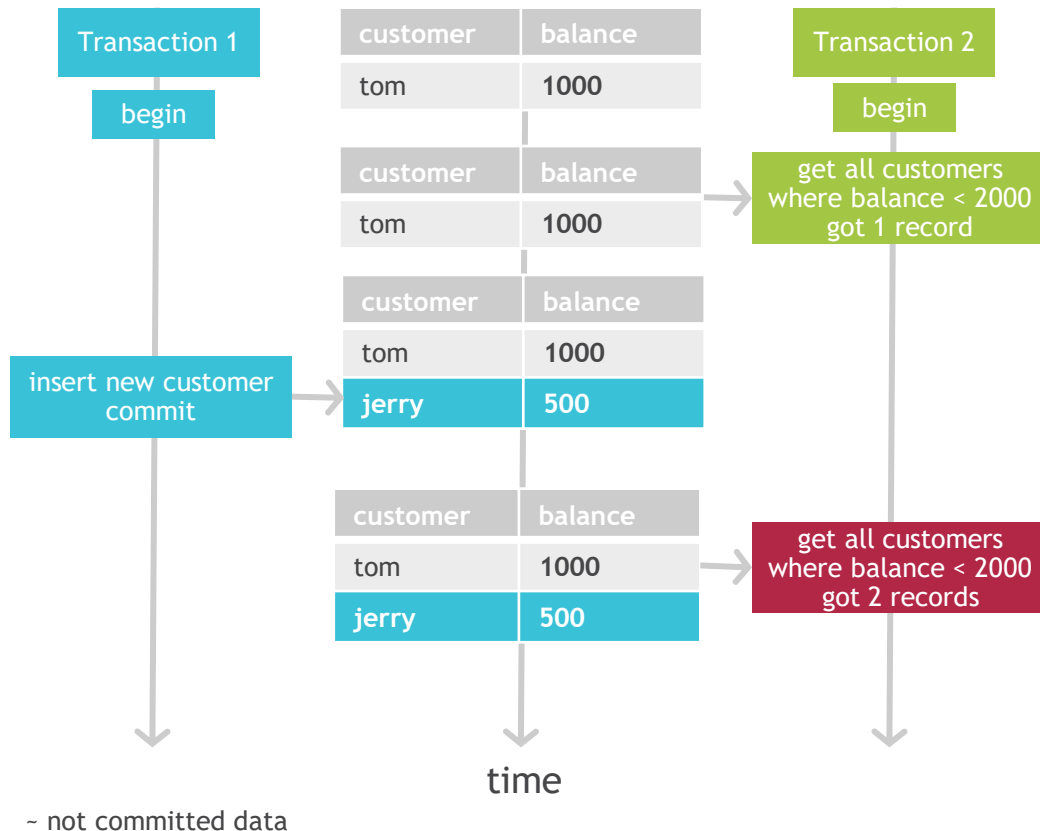
# Transaction phenomena: phantom

## PROBLEM

- One transaction inserts/deletes rows
- Other transaction reads several times and get different number of rows

## WHEN WE CAN LIVE WITH THIS

- Read single rows, not ranges
- We are fine with not the most recent data



HOW TO PROTECT?

# Transaction isolation levels (standard)

	READ UNCOMMITTED	READ COMMITTED	REPEATABLE READ	SERIALIZABLE
Dirty read	YES	NO	NO	NO
Non-repeatable read	YES	YES	NO	NO
Phantom	YES	YES	YES	NO

- Defined in SQL92
- Trade-off between performance, scalability and data protection
- Same work performed with the same inputs may result in different answers, depending on isolation level
- Implementation can be VERY DIFFERNT in different databases



- ISO/IEC 9075:1992
- Information technology -- Database languages -- SQL

<http://www.contrib.andrew.cmu.edu/~shadow/sql/sql1992.txt>

# HOW DOES IT WORK?

**IN DIFFERENT DATABASES**

# Optimistic and pessimistic approaches


## PESSIMISTIC

- Locking rows or ranges
- Like ReadWriteLock/synchronized in Java
- Concurrent transactions wait until lock is released



## OPTIMISTIC

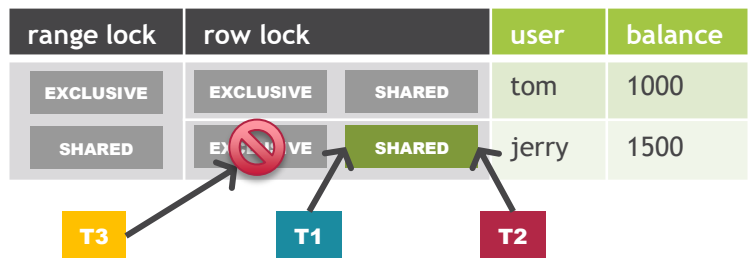
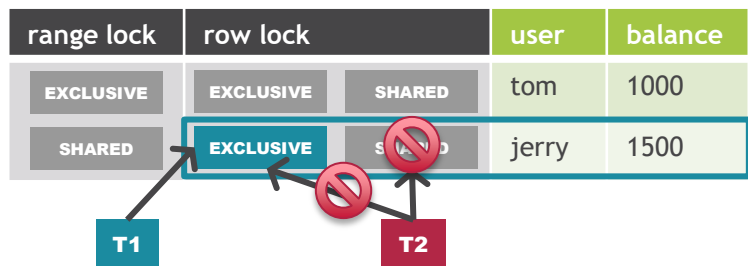
- Multi version concurrency control (MVCC)
- Doesn't lock anything
- Save all versions of data
- We work with data snapshots
- Like GIT/SVN



# Pessimistic locking

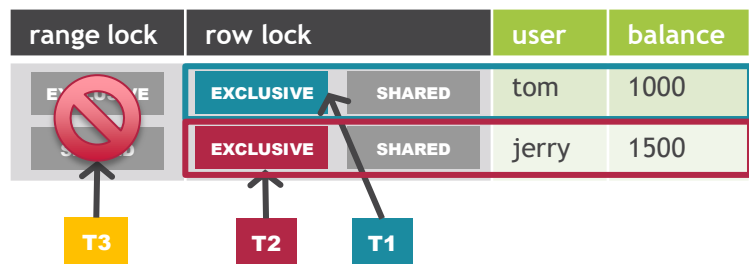

## BY OWNERSHIP

- Shared lock - read lock, many owners
- Exclusive lock - write lock, one owner



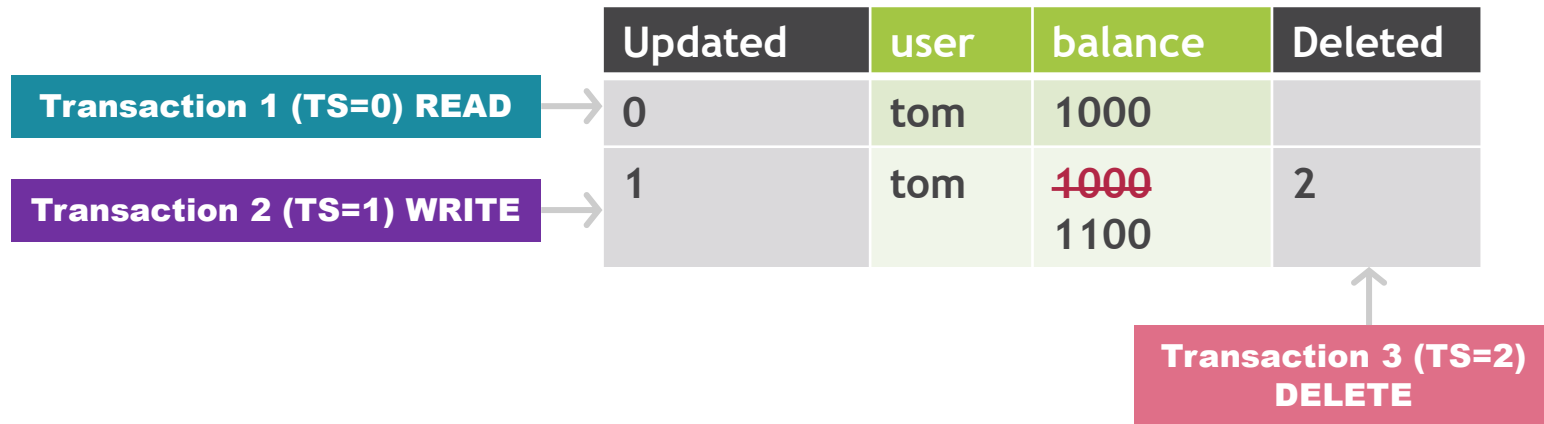
## BY SCOPE

- Row - specific rows by index (if index exists)
- Range - all records by condition





# Optimistic multi version concurrency control (MVCC)



## HOW IT WORKS

- Transactions see the rows with version less or equal to transaction start time
- On update: new version is added
- On remove: deleted column is set of new version number

# Optimistic MVCC vs pessimistic locks

	MVCC (OPTIMISTIC)	LOCKS (PESSIMISTIC)
Behavior	<ol style="list-style-type: none"><li>1. Each transaction works with it's own version</li><li>2. Concurrent transaction fails</li></ol>	<ol style="list-style-type: none"><li>1. Transaction which owns lock works with data</li><li>2. Concurrent transactions wait</li></ol>
Locks	NO	YES
Performance and scalability	GOOD	BAD
Deadlocks	NO	POSSIBLE
Guarantee of recent data version	NO	YES
Extra disk space needed	YES	NO
Durability	better (because of saved versions)	

# Transaction isolation levels in different databases

DB	CONCEPT	READ UNCOMMITTED	READ COMMITTED	REPEATABLE READ	SERIALIZABLE	SPECIFICS
Oracle	MVCC	<b>NOT SUPPORTED</b>	<u>DEFAULT</u> return new snapshot each read	<b>NOT SUPPORTED</b>	returns snapshot of data at beginning of transaction	+ READ ONLY LEVEL transaction only sees data at the moment of start, writes not allowed  always returns snapshots, transaction fail when concurrent update
MySQL (InnoDB)	MVCC		return new snapshot each time	<u>DEFAULT</u> - save snapshot at first read - return it for next reads	- locks ranges - transaction lifetime - Shared lock on select	
MSSQL	LOCKS	+ Double read phenomena: able to read same row twice while it is migrating to another place on disk	<b>SNAPSHOT (optimistic)</b>	- locks rows - transaction lifetime	- locks ranges - transaction lifetime - selects: shared range lock - updates: exclusive lock	+ <b>SNAPSHOT LEVEL</b> - save snapshot at first read - return it for next reads - transactions fail in case of optimistic lock concurrent update
			return new snapshot each time			
			<b>LOCK (pessimistic)</b>			
			<u>DEFAULT</u> - locks rows - statements lifetime			
PostgreSQL	MVCC	<b>NOT SUPPORTED</b>	<u>DEFAULT</u> return new snapshot each read	- save snapshot at first read - return it for next reads	predicate locking (optimistic)	always returns snapshots, transaction fail when concurrent update

**TRANSACTION ISOLATION IN MYSQL**

**LIVE DEMO**

**WANT TO OPTIMIZE?**

**LOCK SPECIFIC OBJECTS!**

# Pessimistic locking of specific rows/ranges (MySQL)

## LOCKING SELECTS

- `SELECT ... LOCK IN SHARE MODE` - shared (read) lock
- `SELECT ... FOR UPDATE` - exclusive (write) lock


## IDEA

- Increase isolation level for specific rows/ranges
- Other rows/ranges can have lower isolation level

**MYSQL PESSIMISTIC LOCKING OF SPECIFIC OBJECTS**

**LIVE DEMO**

# Database deadlocks



Database deadlocks happen because of bad application architecture design

## HOW TO PREVENT DEADLOCKS

- Take locks in same order in every transaction
- Use as small as possible transactions



WHAT ABOUT JAVA?

# Transaction isolation: configuration in Java

## JDBC

```
Connection c = DriverManager.getConnection("jdbc:mysql://localhost/testdb?user=test&password=pass");  
c.setTransactionIsolation(level);
```

## Hibernate

```
hibernate.connection.isolation=level
```

Connection.TRANSACTION_NONE	0
Connection.TRANSACTION_READ_UNCOMMITTED	1
Connection.TRANSACTION_READ_COMMITTED	2
Connection.TRANSACTION_REPEATABLE_READ	4
Connection.TRANSACTION_SERIALIZABLE	8

# JPA features for locking

## Enum LockModeType

PESSIMISTIC\_READ

Shared lock

PESSIMISTIC\_WRITE

Exclusive lock

## EntityManager

`lock(Object entity, LockModeType lockMode)`

Makes additional locking select query just to lock entity

`find(Class<T> entityClass, Object primaryKey, LockModeType lockMode)`

Makes locking select when reading entity

`refresh(Object entity, LockModeType lockMode)`

Makes locking select when reloading entity

## NamedQuery

`@NamedQuery(name="myQuery", query="...", lockMode=LockModeType.PESSIMISTIC_READ)`

Allows to specify that we need locking select to any query

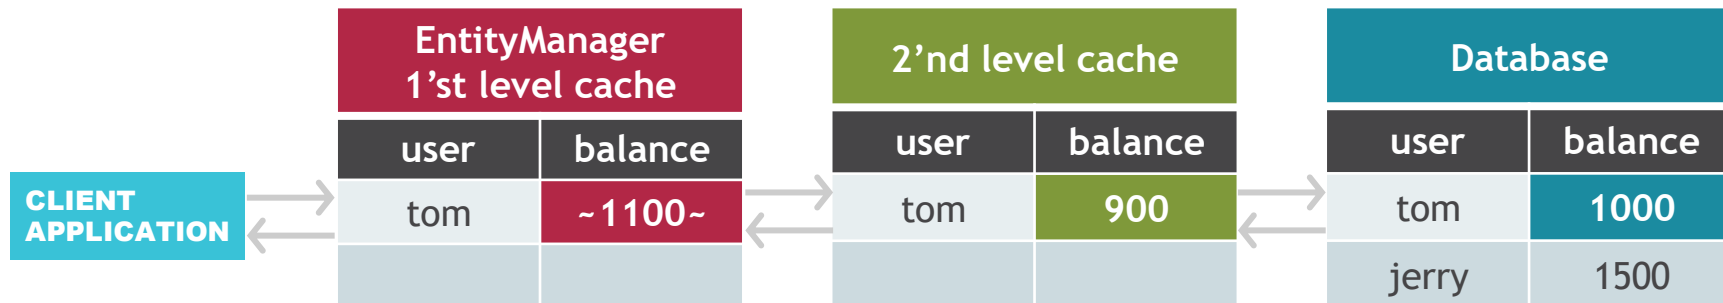
## ADVANTAGES

- It is really simple
- Database specific things are hidden from developer
- Supports parent-child entities

## DRAWBACKS

- Complex to manually lock entity relationships
- `@NamedQuery` is only way to specify query lock

# Transaction isolation and locking with JPA




- Repeatable reads because of EntityManager's cache
- Requests do not always go to database

Behavior = EntityManager + 2'nd lvl cache + database


# Conclusion

- Do you have problems because of concurrent updates?
  - Same as concurrent programming in Java
  - Sometimes we can allow phenomena
- Transaction Isolation is trade-off between data protection and performance
- Two main approaches in databases implementation:
  - Optimistic: no locks, data is versioned
  - Pessimistic: range and low locks
- JPA
  - Simplifies usage of pessimistic locking
  - Adds own specific behavior because of caches
- For better performance:
  - Prefer smaller transactions: they hold long time locks and can make deadlocks
  - Be careful with declarative transaction management it can make heavy transactions


# References




Oracle  
Chapter 2-4




MySQL  
Chapter 1



JPA  
Chapter 12,  
Locking



MSSQL  
Chapter 13



PostgreSQL  
Chapter 10



Examples  
[https://github.com/  
kslisenko/tx-isolation](https://github.com/kslisenko/tx-isolation)



THANK YOU!

QUESTIONS?

kanstantsin\_slisenka@epam.com