



Networking in Java with NIO and Netty

KANSTANTSIN SLISENKA
SENIOR SOFTWARE ENGINEER

DECEMBER 22, 2016

About me



Kanstantsin Slisenka

- Java backend developer
- Speaker at Java tech talks

I AM INTERESTED IN

- Complex Java backend, SOA, databases
- High load, fault-tolerant, distributed systems

skype: kslisenko

kslisenko@gmail.com, kanstantsin_slisenka@epam.com

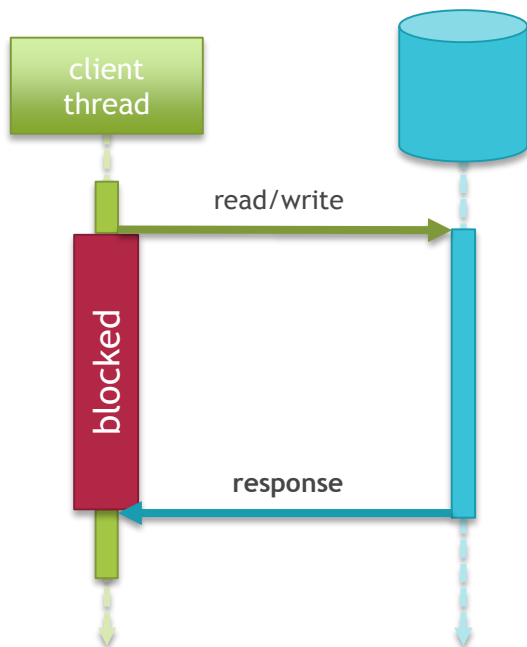
Agenda

- 1 Blocking vs non-blocking resources
- 2 Blocking IO: sockets
- 3 Non-blocking IO: NIO socket channels
- 4 Netty: high performance networking framework

Blocking vs non-blocking resources

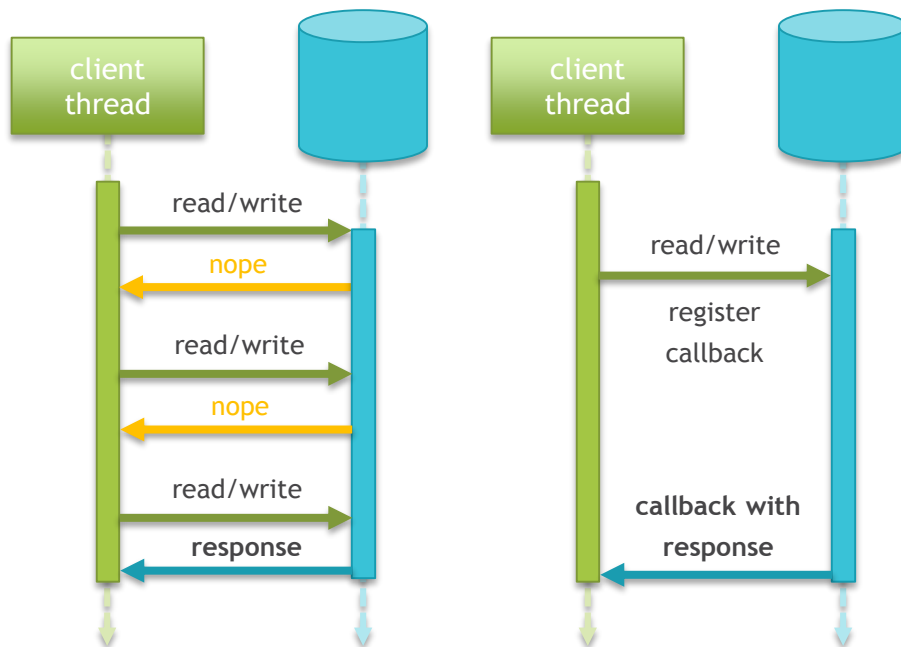
Blocking resources

Database transactions, RPC/WS calls

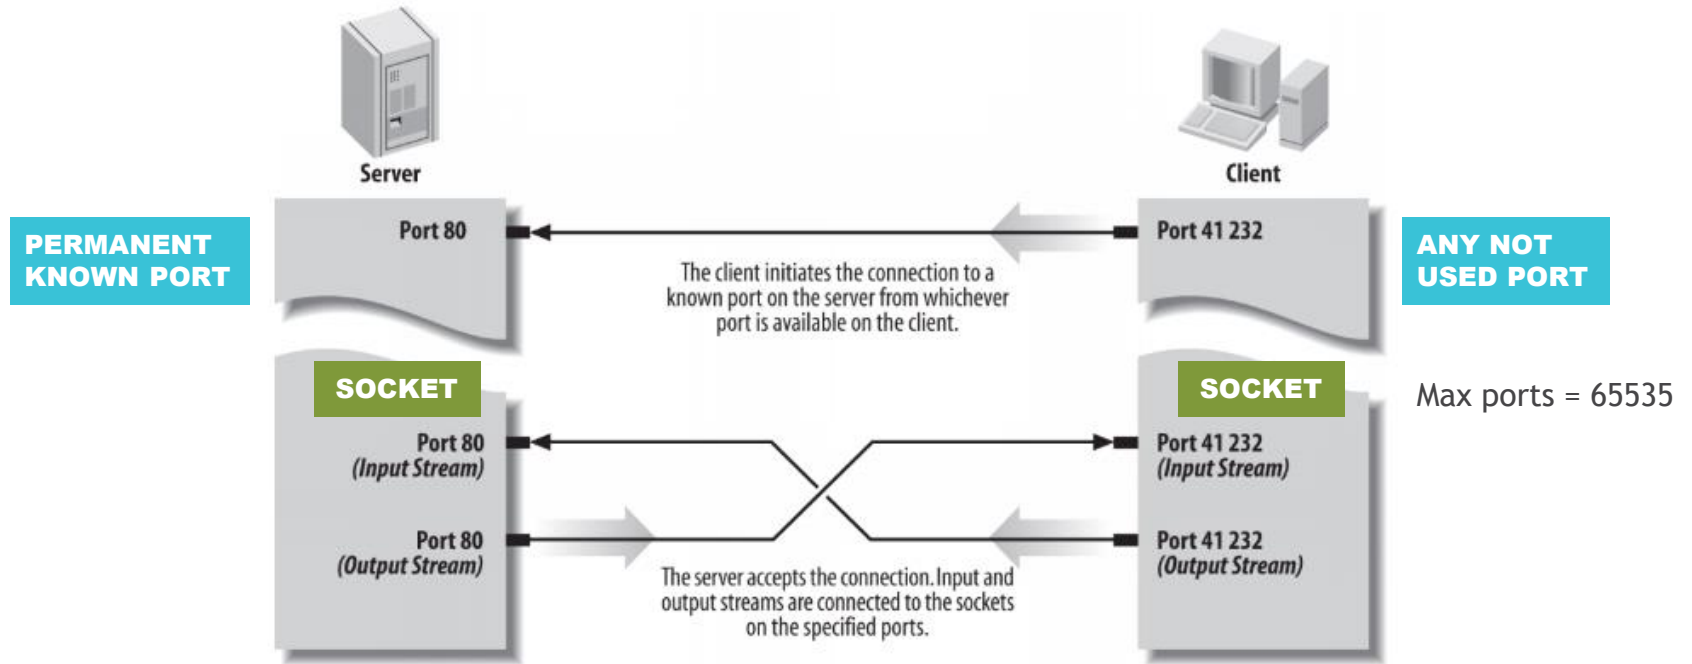


Non-blocking resources

Queues, WebSockets, Non-blocking data storage



Network sockets



Sockets are peer-to-peer communication objects at both sides

Java sockets use blocking IO

Client

```
Socket socket = new Socket("localhost", 8080);
OutputStream out = socket.getOutputStream();
InputStream in = socket.getInputStream();

out.write("hello".getBytes()); // Blocking call

byte[] data = new byte[1024];
while (in.read(data) != -1) { // Blocking call
    handle(data);
}
```

Server

```
ServerSocket serverSocket = new ServerSocket(8080);
Socket socket = serverSocket.accept();
InputStream in = socket.getInputStream();
OutputStream out = socket.getOutputStream();

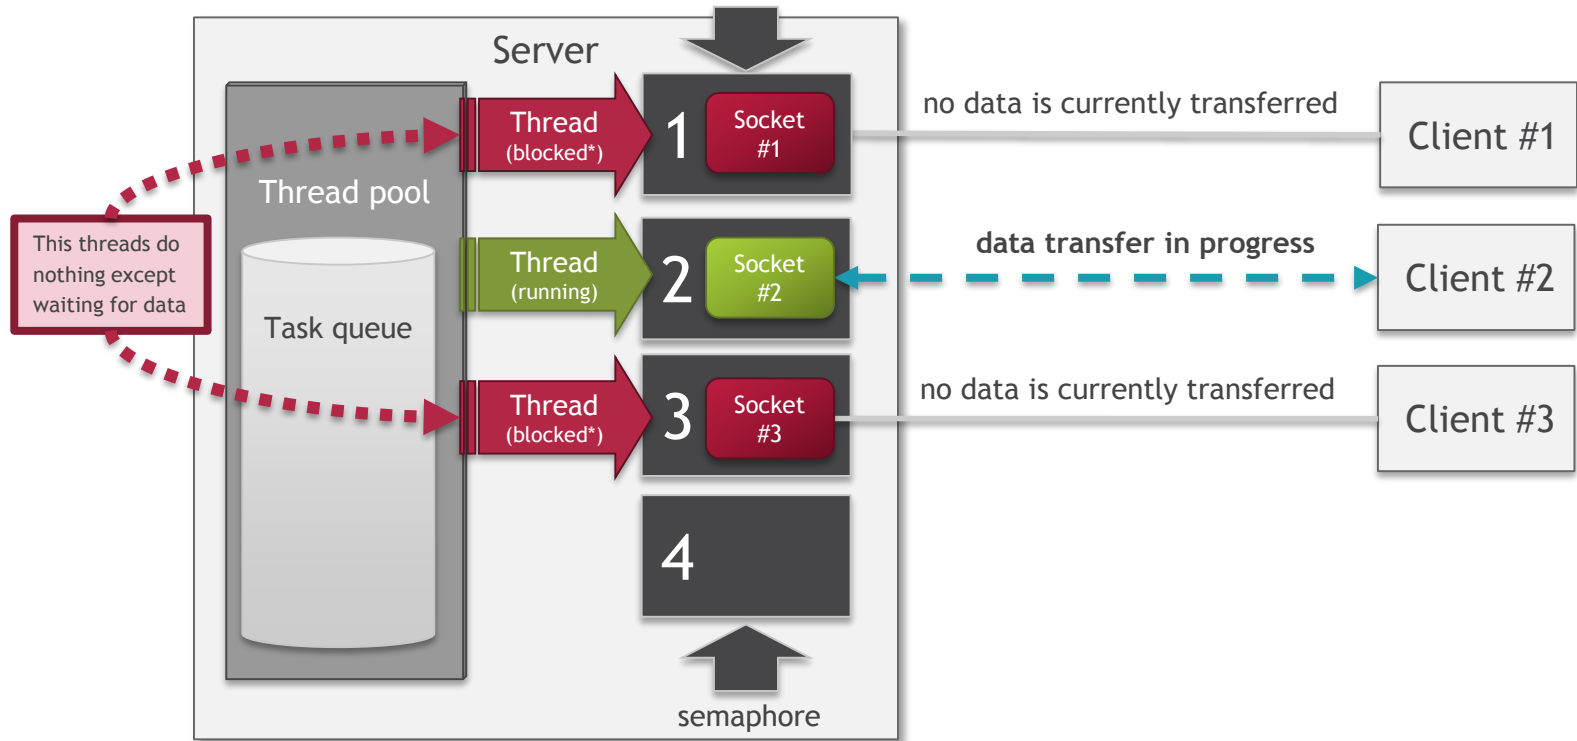
byte[] data = new byte[1024];
while (in.read(data) != -1) { // Blocking call
    out.write(data); // Blocking call
}
```


A close-up, high-angle shot of a person's hands typing on a silver laptop keyboard. The person is wearing a green and white plaid shirt. The laptop is open, and a pair of glasses is resting on the desk next to it. A teal text overlay is positioned in the lower-left quadrant of the image.

LIVE CODING EXAMPLE

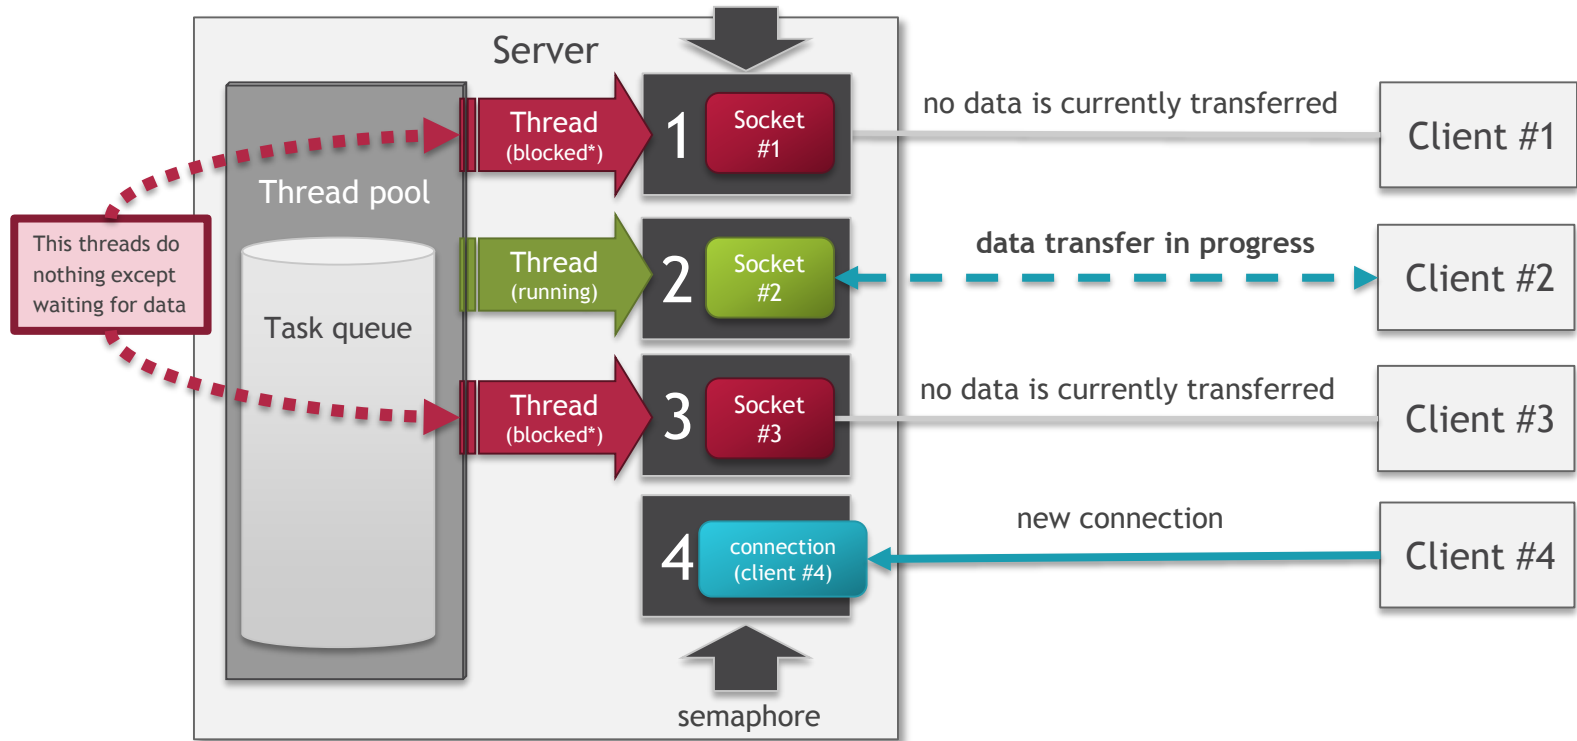
SOCKET SERVER

Socket server architecture



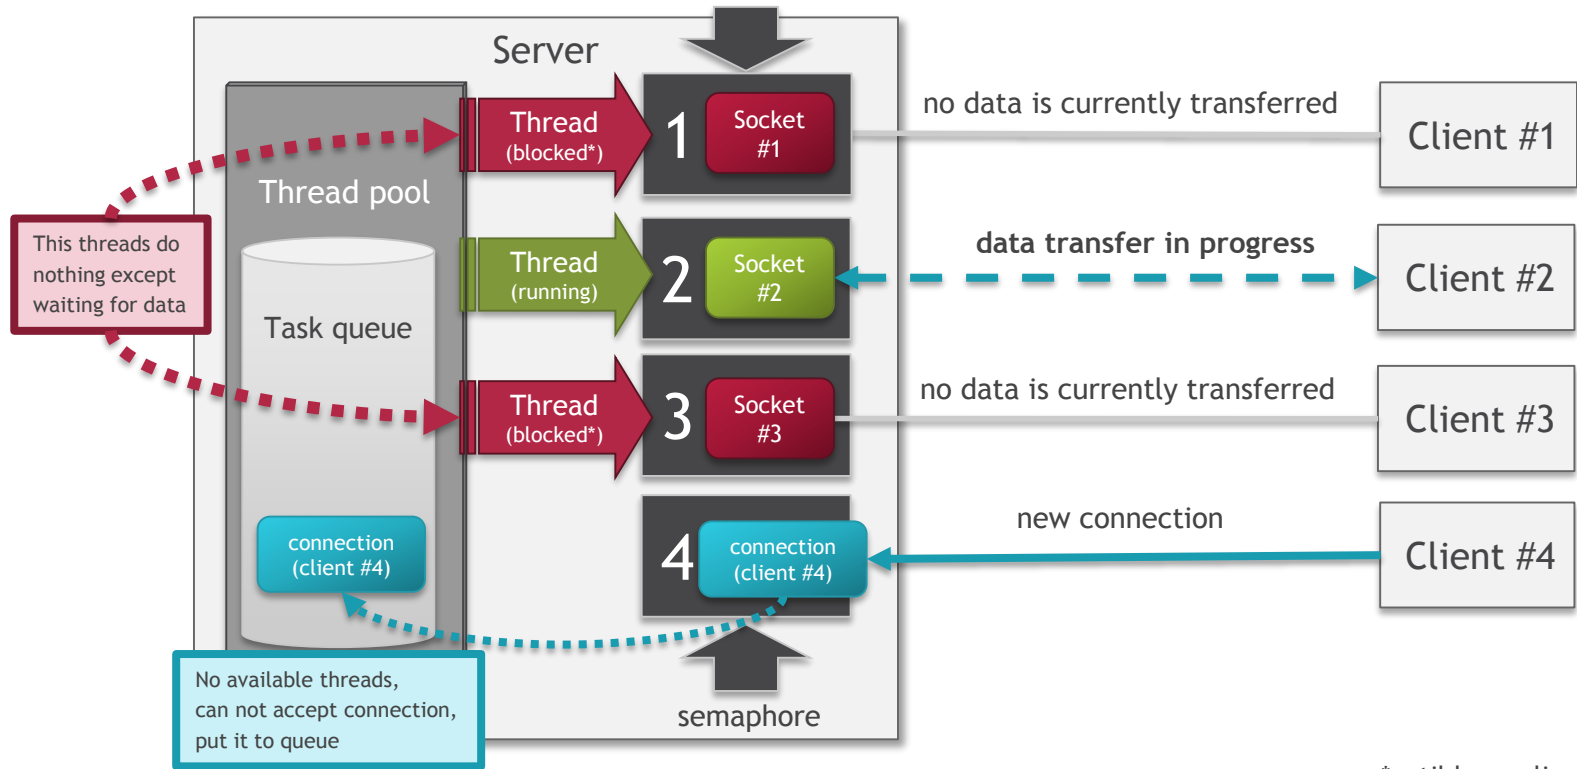
*until keep alive timeout

Socket server architecture

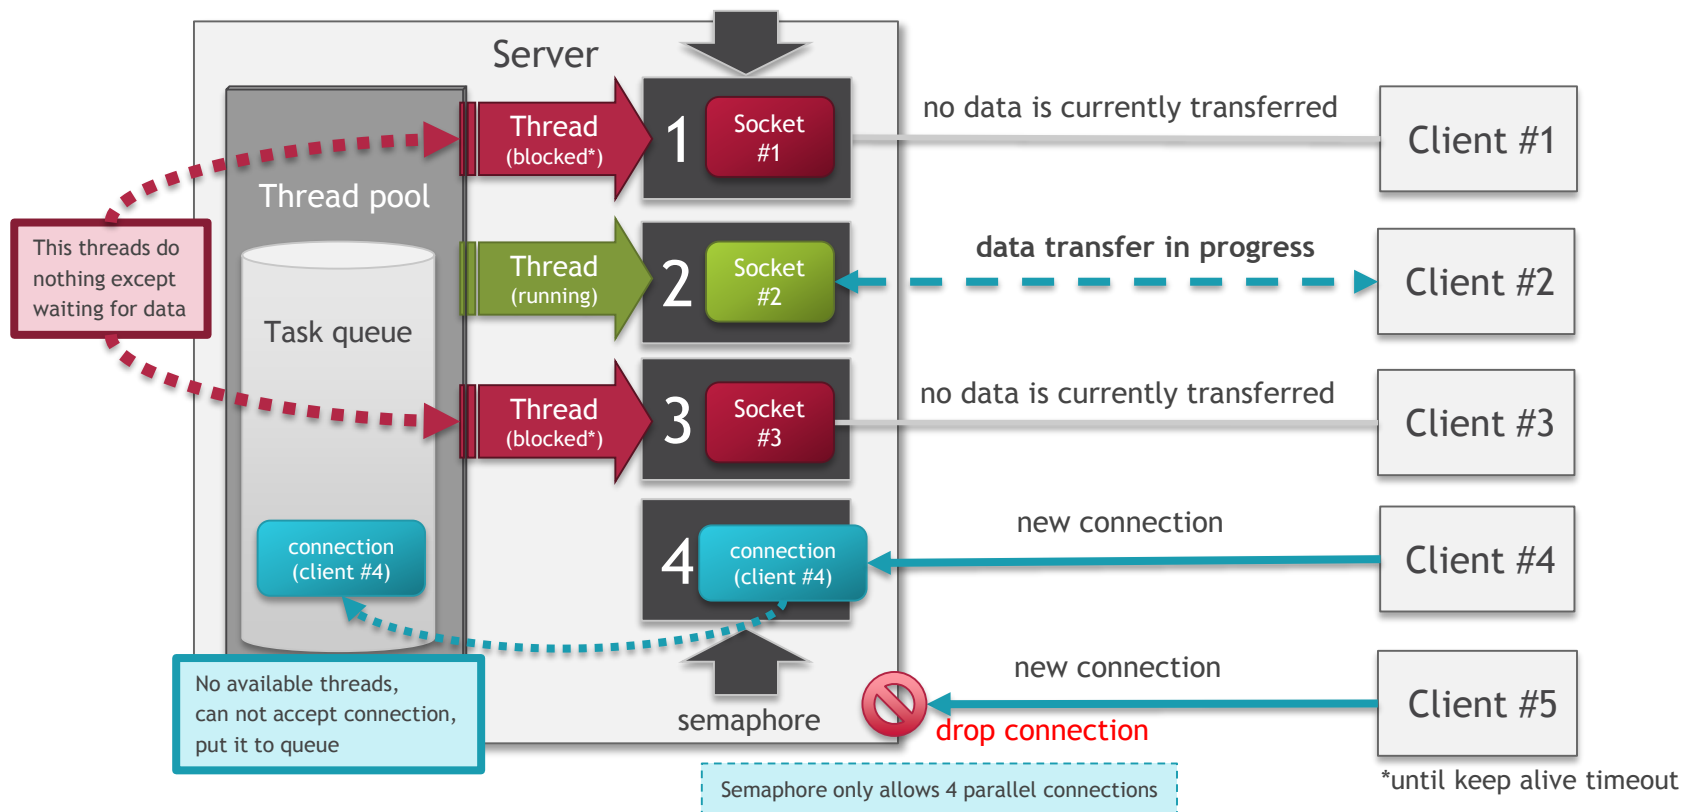


*until keep alive timeout

Socket server architecture



Socket server architecture



Socket server conclusion

Advantages

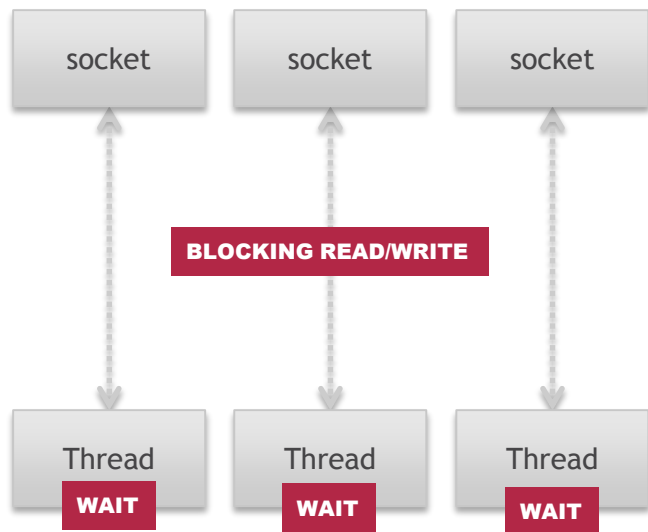
- **Simple development**
because we are using well known
Input/OutputStream API
- **No check for partially received messages**
thread is blocked until fully reads message

Drawbacks

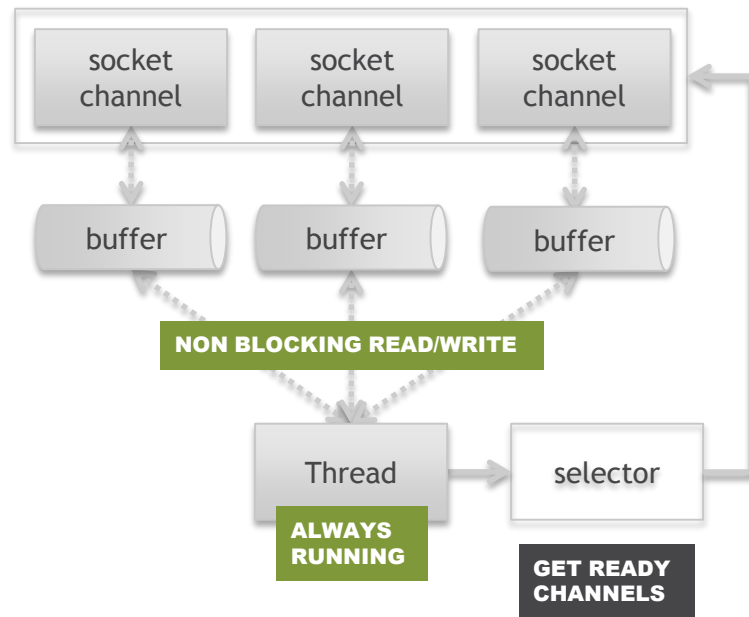
- **Too many threads**
extra memory
context switching overhead
- **Responsiveness**
when all threads are blocked putting new
connections to the thread pool queue makes
server unresponsive
- **Connections count limit**
semaphore prevents from over queuing, but
limits number of parallel connections

IO vs NIO

IO (blocking)




NIO (non-blocking)



NIO byte buffer


Empty buffer



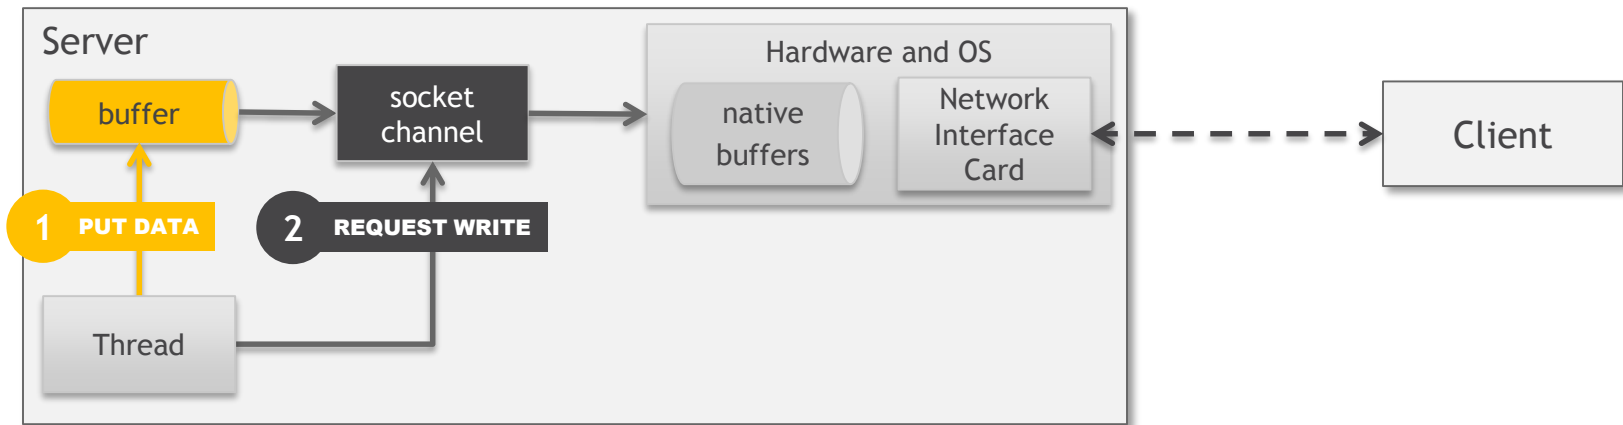
Writing mode

`buffer.clear()` `buffer.flip()`

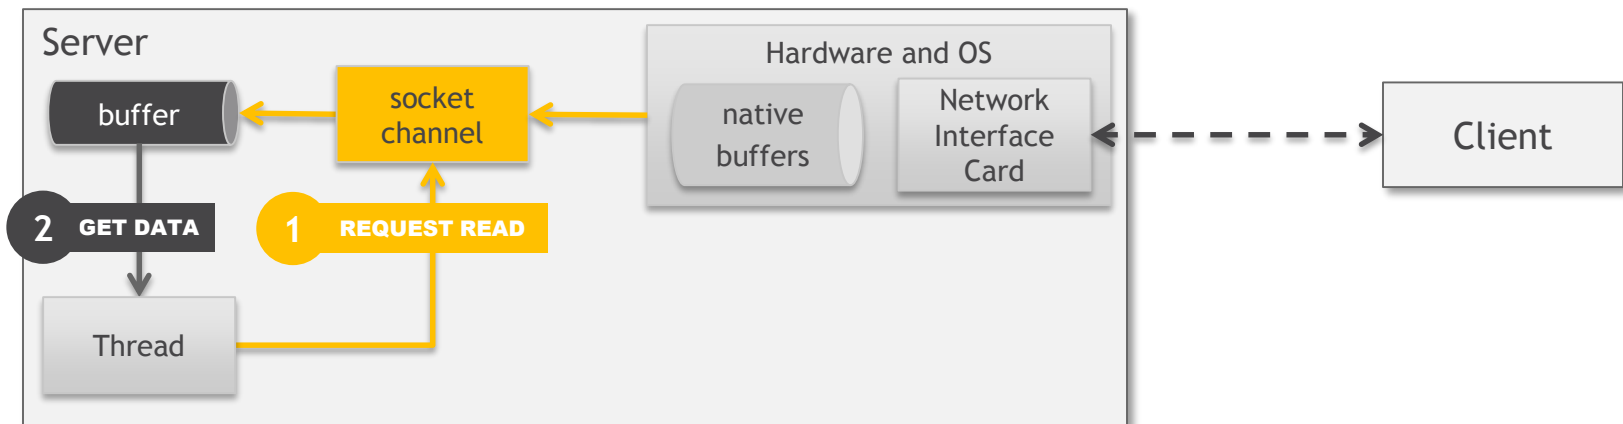
Reading mode



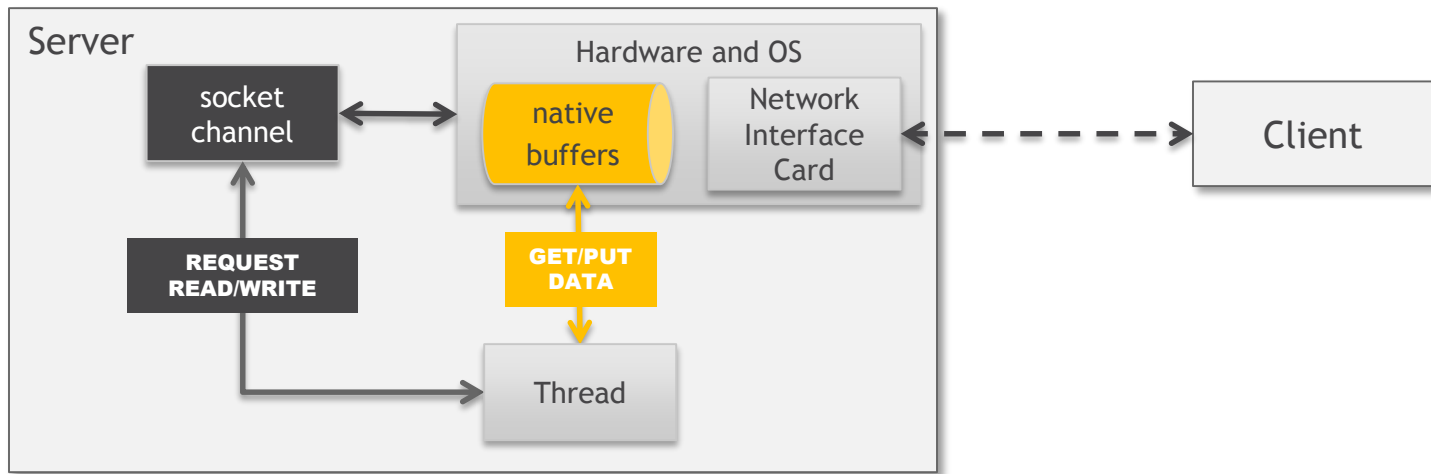
NIO socket channel: data sending



NIO socket channel: data receiving



NIO direct buffer



Lives in native memory

Cheap to read/write

Expensive to allocate/deallocate

NIO selector



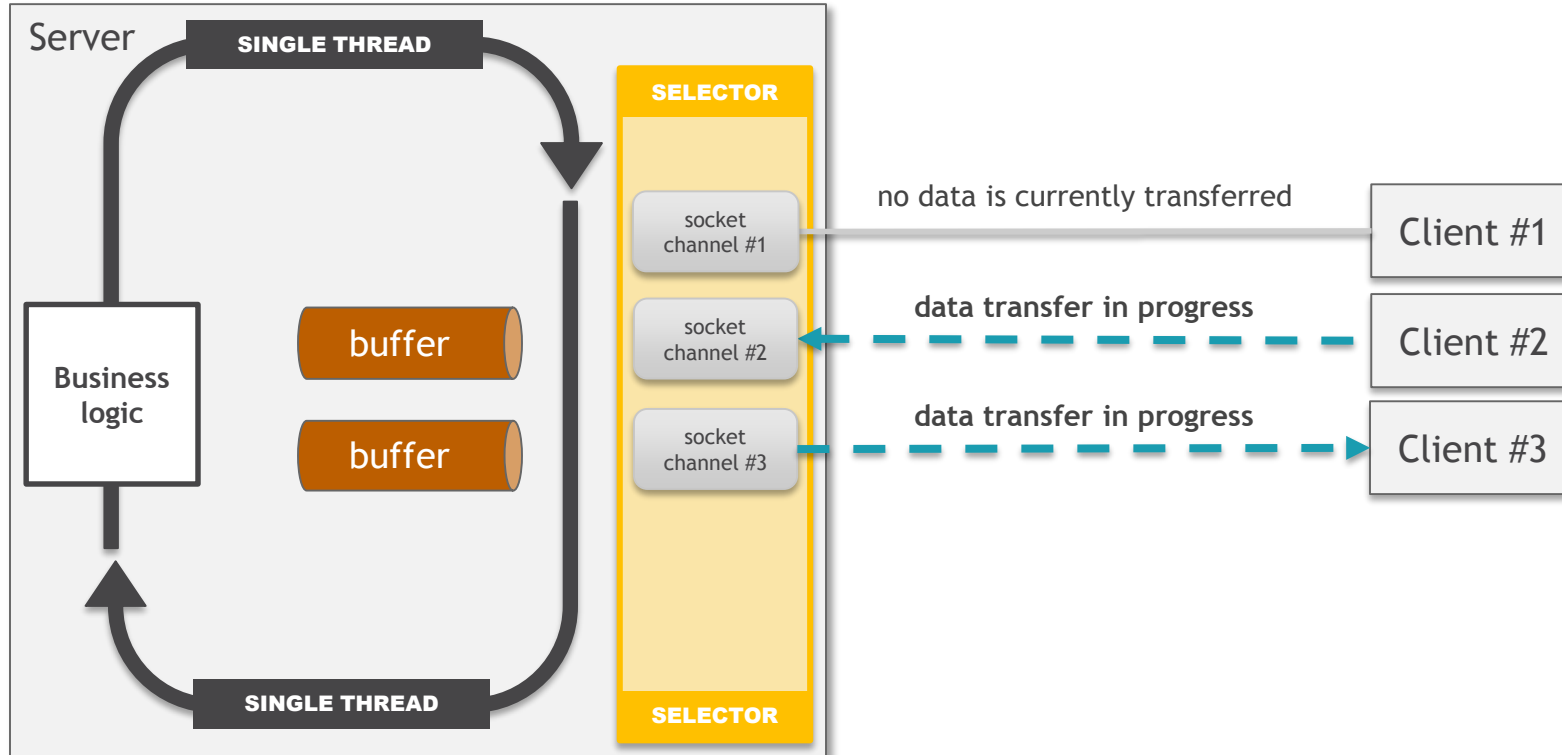
Selector blocks until something happens at least at one channel and returns ready channels



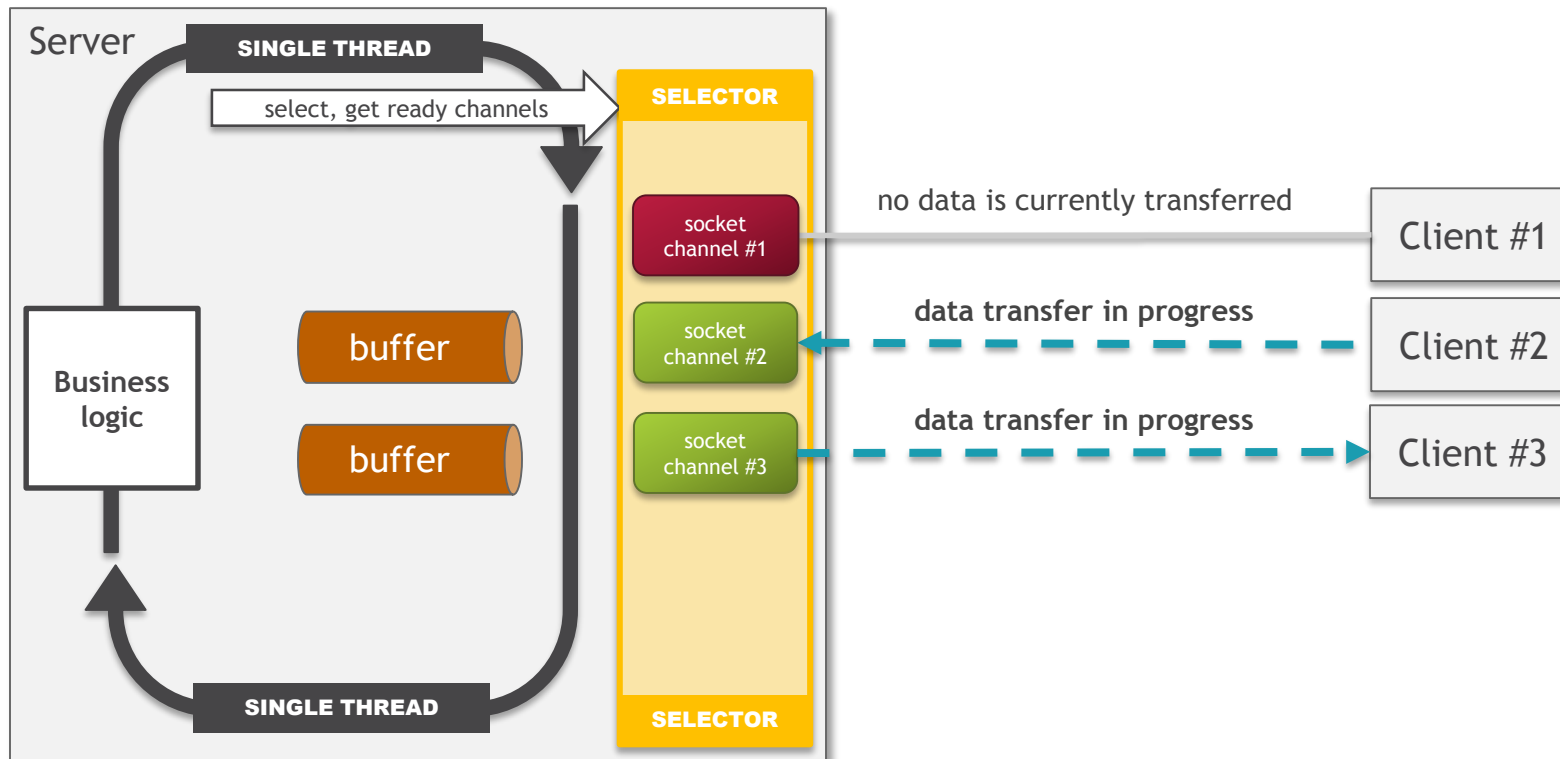
LIVE CODING EXAMPLE

NIO SERVER

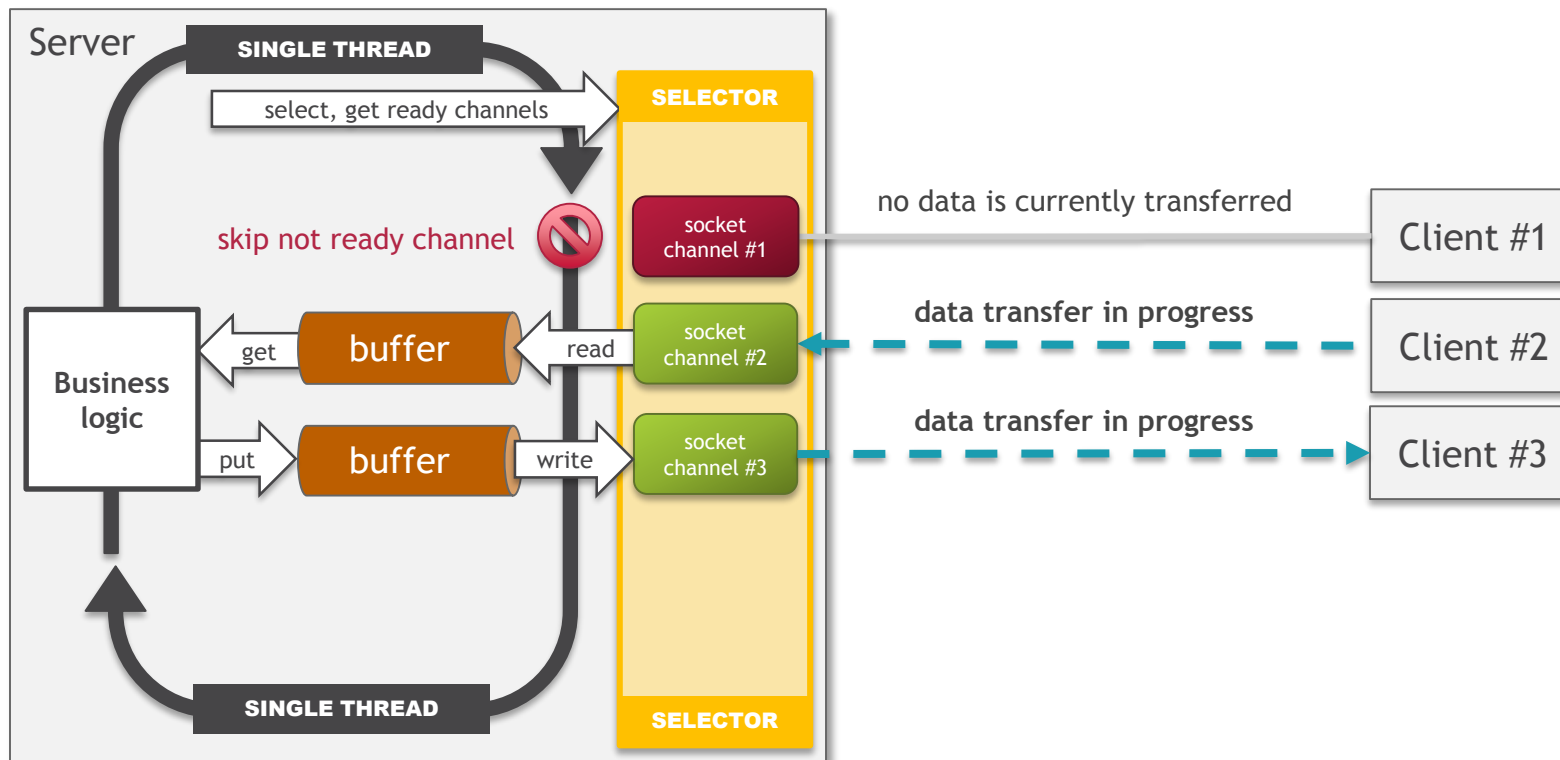
Basic NIO server architecture



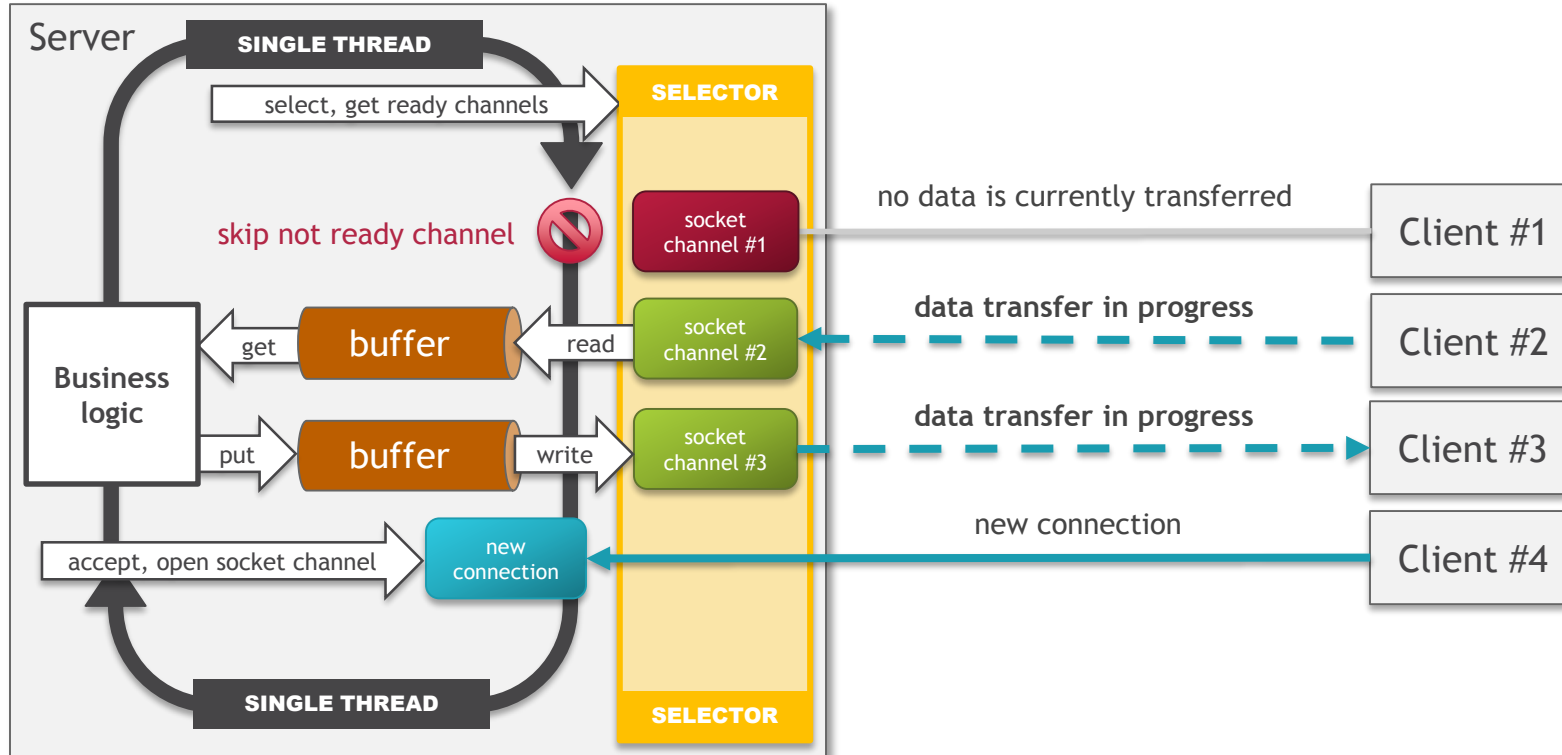
Basic NIO server architecture



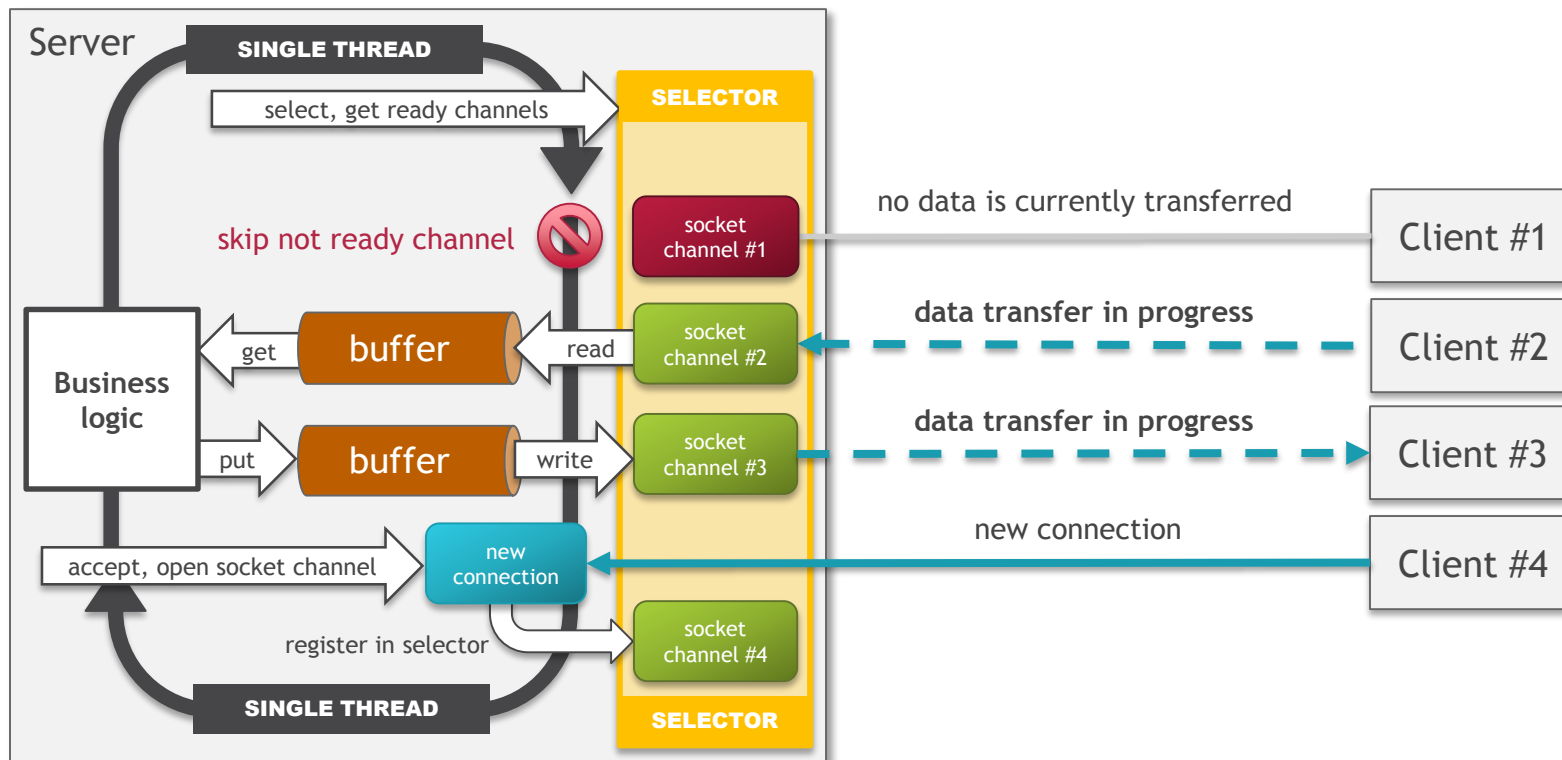
Basic NIO server architecture



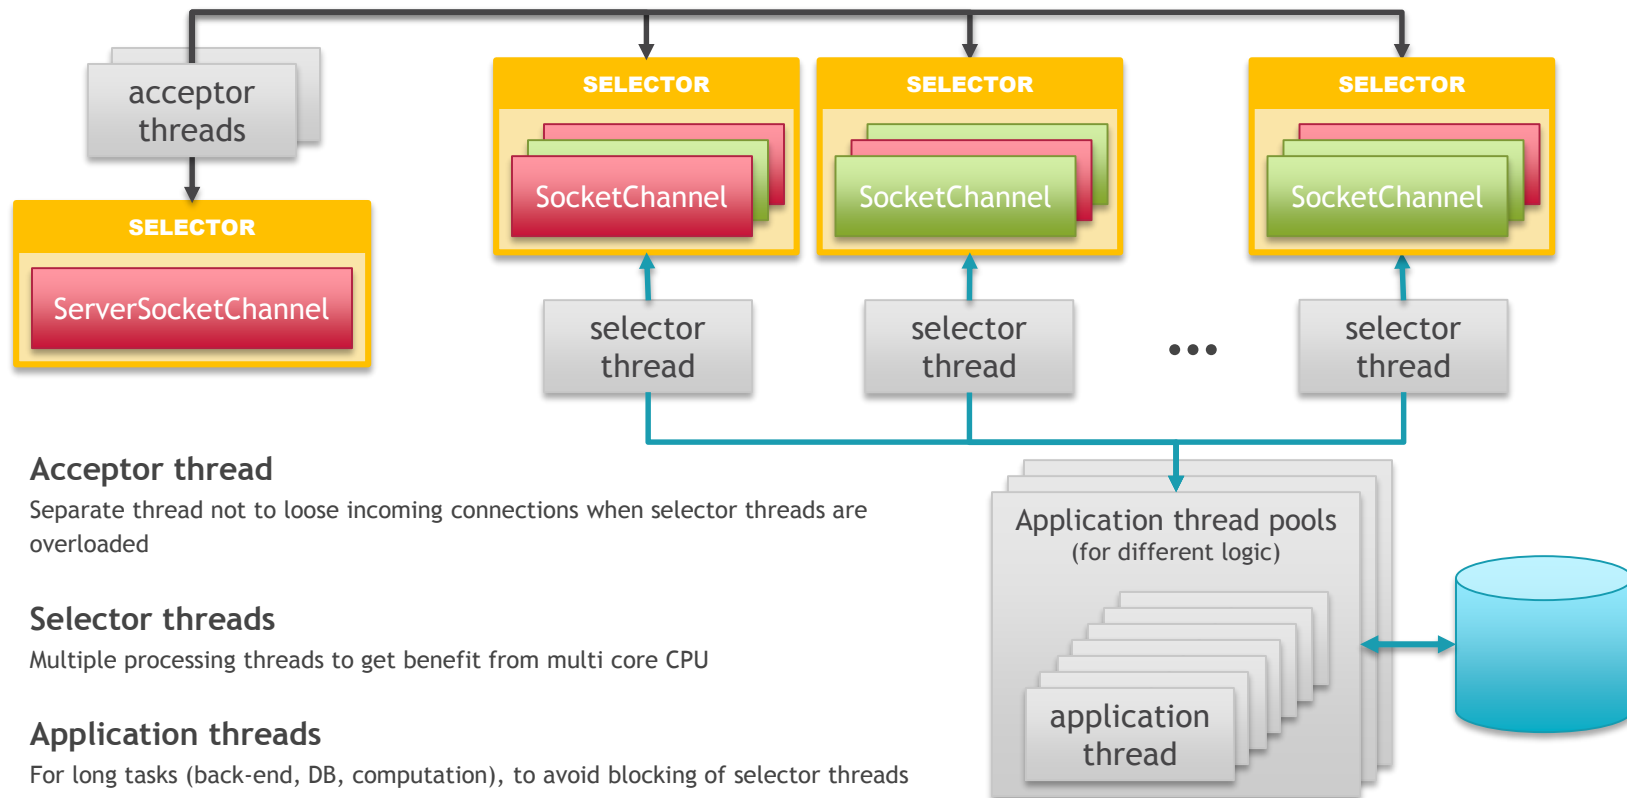
Basic NIO server architecture



Basic NIO server architecture




Production-ready NIO server



Partial message receiving problem

Message reader splits/joins data blocks into messages



NIO conclusion

Advantages

Single thread for many connections

- less memory
- less context switching overhead

Drawbacks

Complexity

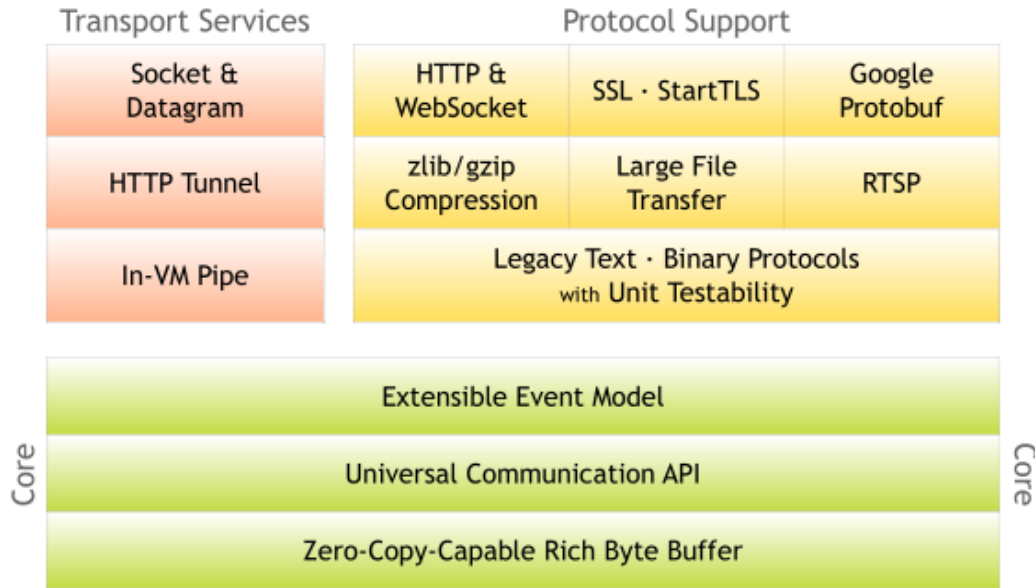
- complex code
- handling partially received messages
- pain when we combine asynchronous and synchronous code - need different thread groups

A cyclist wearing a black jersey, shorts, and a helmet is riding away from the viewer on a rocky, uneven path. In the background, a large, conical volcano rises against a blue sky with scattered white clouds. The foreground is filled with dark, jagged rocks and patches of green grass with small white daisies.


THE WAY FROM PAIN TO PLEASURE

NETTY: ONE FRAMEWORK TO RULE THEM ALL

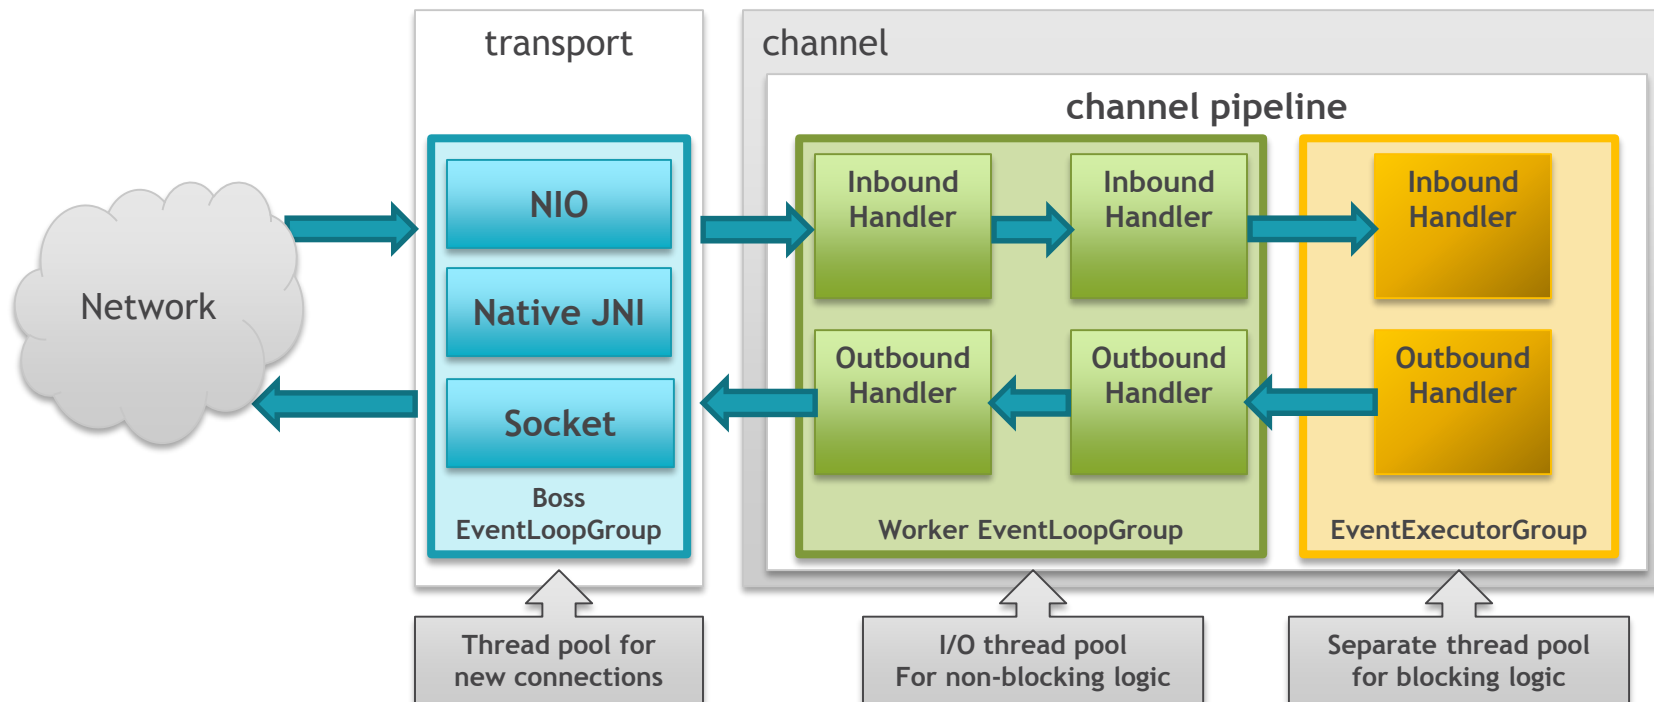
Netty: network application framework



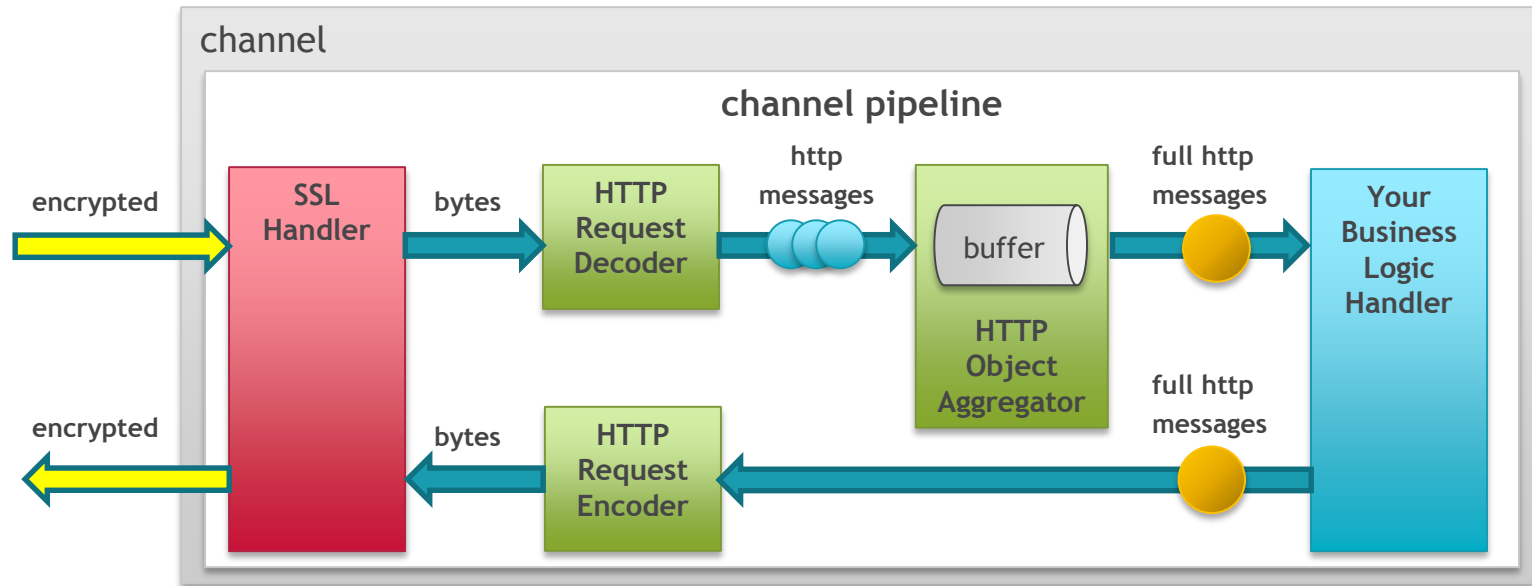
Some projects using Netty



Netty server architecture



Netty handlers example




A close-up photograph of a person's hands typing on a silver laptop keyboard. The person is wearing a green and white plaid shirt. The laptop is open, and a pair of glasses is resting on the desk next to it. A teal text overlay is positioned in the lower-left quadrant of the image.

LIVE CODING EXAMPLE

NETTY SERVER


What is faster: IO or NIO?



What is faster: IO or NIO?

May be, may be not *

1. ... NIO is faster
2. ... thread context switching is slow
3. ... threads take up too much memory
4. ... synchronization among huge number of threads will kill you
5. ... thread per connection does not scale



* Always point to holy-wars, see

www.mailinator.com/tymaPaulMultithreaded.pdf

What is faster: IO or NIO?

May be, may be not *

1. ... NIO is faster
Benchmark shows IO 25-30% faster*
2. ... thread context switching is slow
New threading library in Linux 2.6
(en.wikipedia.org/wiki/Native_POSIX_Thread_Library)
Idle thread cost is near zero
Context switching is much much faster
3. ... threads take up too much memory
Yes, it is, Linux (x64) default thread
stack size (xss) = 1024kb
4. ... synchronization among huge number of
threads will kill you
Non-blocking data structures which scale well
5. ... thread per connection does not scale
See 2-4



* Always point to holy-wars, see

www.mailinator.com/tymaPaulMultithreaded.pdf

Conclusion

Use sockets*


- Small number of connections
- Huge data rate per connection
- Synchronous world (procedure-based)
 - Classic enterprise applications
 - Distributed XA transactions
 - Request-response protocols
 - Blocking database
 - Blocking RPC or web-service calls
- Bottleneck is not server, but blocking resource
- Make servers stateless and scale horizontally

Use NIO*

- Huge number of connections
- Medium or slow data rate per connection
- Asynchronous world (event-based)
 - Non blocking database
 - Queues
 - Web-sockets
 - Callbacks, listeners
 - Data streaming
- Don't use standard NIO, use frameworks like Netty
- Highly skilled team because of complex code


**Not a strict rule, always point to holy-wars*

BUY A BOOK, MAKE THEM RICH




Java Network Programming, 4'th edition

O'REILLY
Elliotte Rusty Harold



Java NIO

O'REILLY
Ron Hitchens



Netty in Action

Manning
Norman Maurer

- **Code examples:** github.com/kslisenko/java-networking
- **NIO performance:** www.mailinator.com/tymaPaulMultithreaded.pdf
- **NIO tutorial:** tutorials.jenkov.com/java-nio/index.html
- **ITT 2015 - Heinz Kabutz - The Multi-threading, Non Blocking IO:** youtube.com/watch?v=uKcOGx_lPsg
- **Netty - One Framework to rule them all by Norman Maurer:** youtube.com/watch?v=DKJ0w30M0vg
- **Scalable IO in Java:** <http://gee.cs.oswego.edu/dl/cpjslides/nio.pdf>

THANK YOU AND HAPPY HOLIDAYS!

QUESTIONS?

KANSTANTSIN_SLISENKA@EPAM.COM